

COMP 3011
DESIGN AND ANALYSIS OF ALGORITHMS
FALL 2024

Graphs and Greedy Algorithms

LI Bo
Department of Computing
The Hong Kong Polytechnic University

COMMENTS

Comment: I think for the difficulties part, you can provide some solid **supplementary sources** on BB for those who never learn anything about algorithms. And then **keep the difficulties** in lecture for the advanced students.

- Solid supplementary sources on BB
 1. VIII Appendix: Mathematical Background (Textbook: Introduction to Algorithms)
 2. Self-Reading (Math Foundations) on BB
- Keep the difficulties in lecture for the advanced students
 1. Will Keep the difficulty standard.
 2. Look at the materials from the angles of design and analysis.

COMMENTS

- We have tutorial next Monday (Sep 16, 2024)
- We do not have lectures next Wednesday (Sep 18, 2024)

Happy Mooncake festival!

BIPARTITE GRAPHS

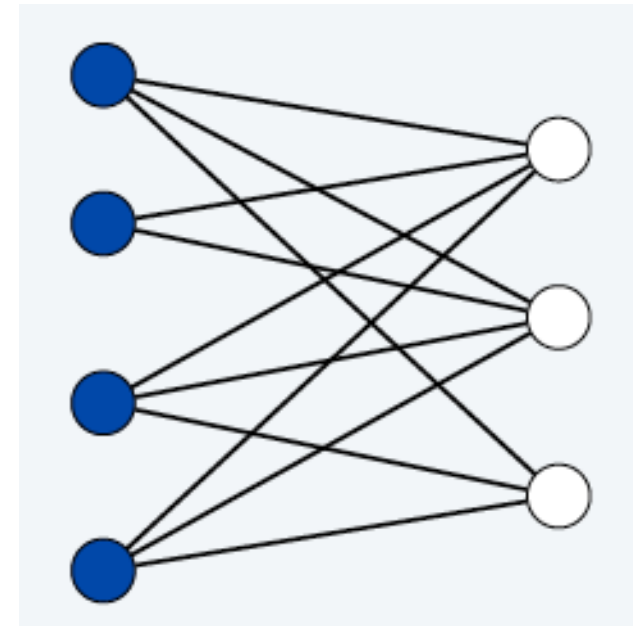
BIPARTITE GRAPHS

Definition (Bipartite Graphs)

- An undirected graph $G = (V, E)$ is **bipartite** if the nodes can be coloured **blue or white** such that every edge has one white and one blue end.

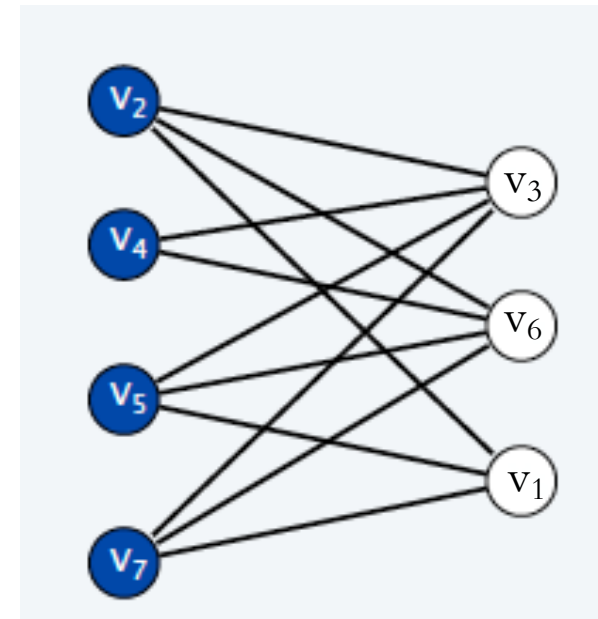
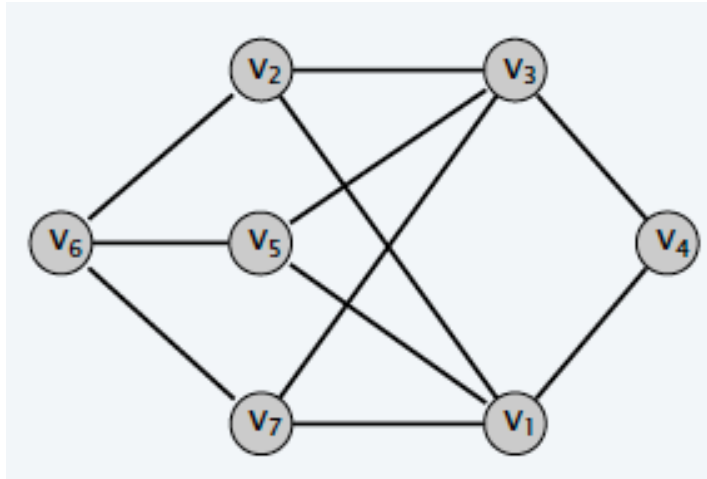
Application

- **Stable matching**: med-school students = blue,
hospitals = white.



BIPARTITE GRAPHS

- Is this graph bipartite?



Given a graph, how to
determine whether it is
bipartite or not?

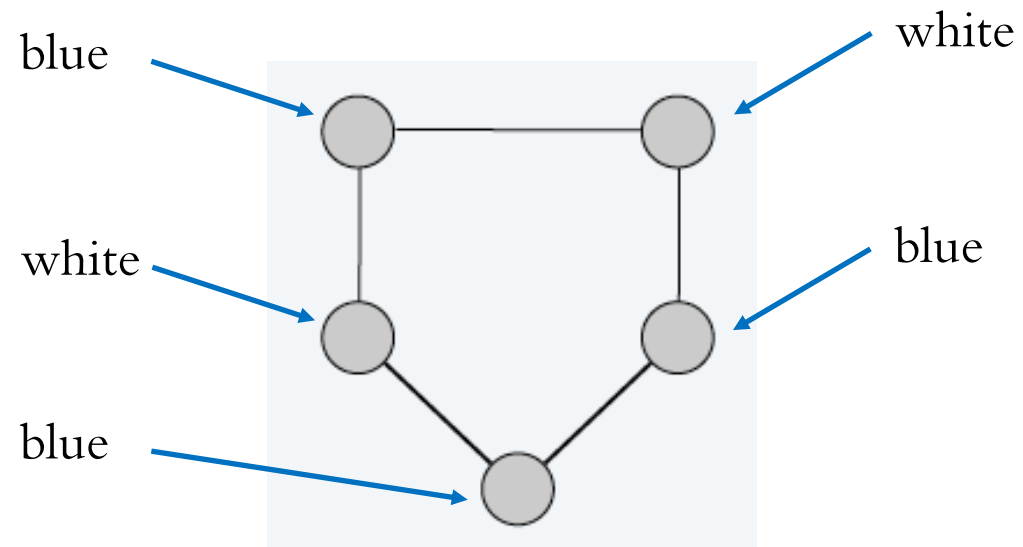
BIPARTITE GRAPHS

An edge has different colors on the two end points.

Lemma. If a graph G is bipartite, it cannot contain an odd-length cycle.

Proof:

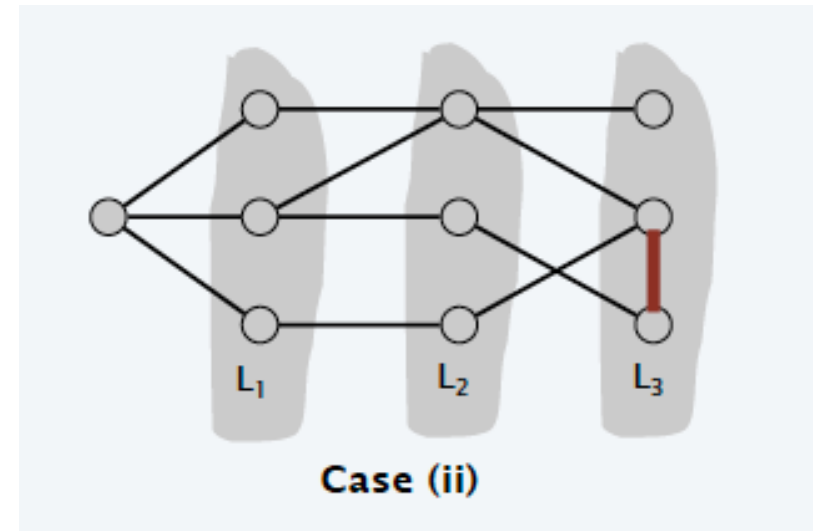
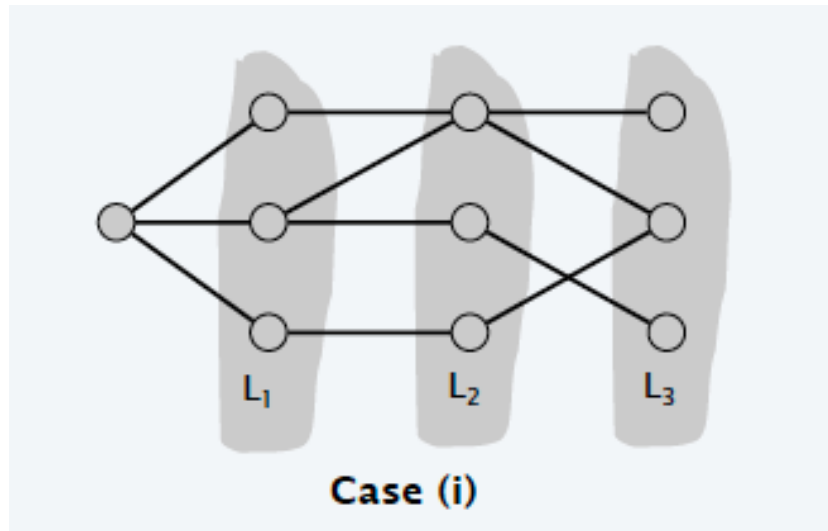
➤ Not possible to 2-color the odd-length cycle, let alone G . ■



BIPARTITE GRAPHS

Lemma. Let G be a connected graph, and let L_0, \dots, L_k be the layers produced by BFS starting at node s . **Exactly** one of the following holds.

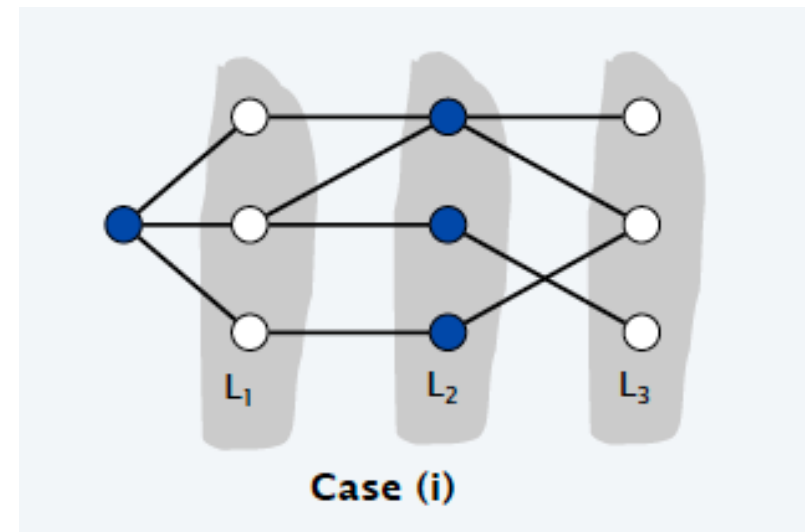
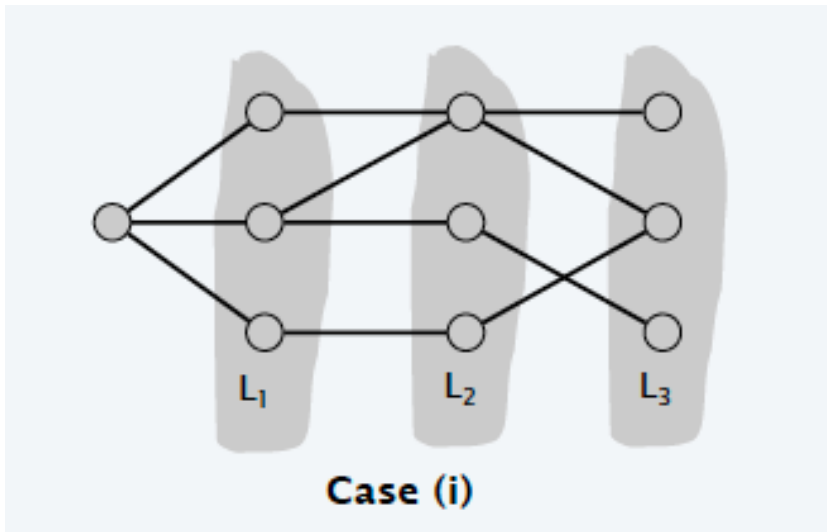
- i. No edge of G joins two nodes of the same layer, and G is **bipartite**.
- ii. An edge of G joins two nodes of the same layer, and G contains an odd-length cycle (and hence is **not bipartite**).



BIPARTITE GRAPHS

Case i:

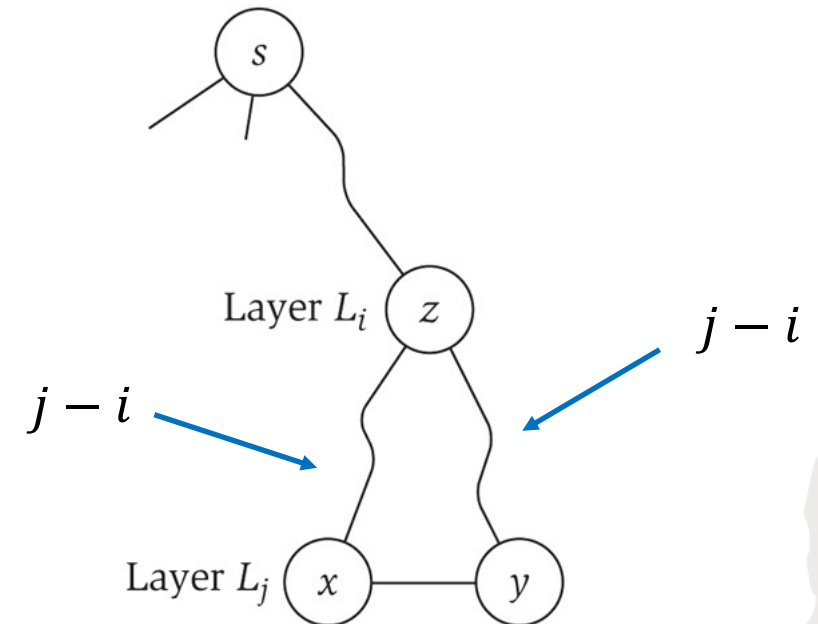
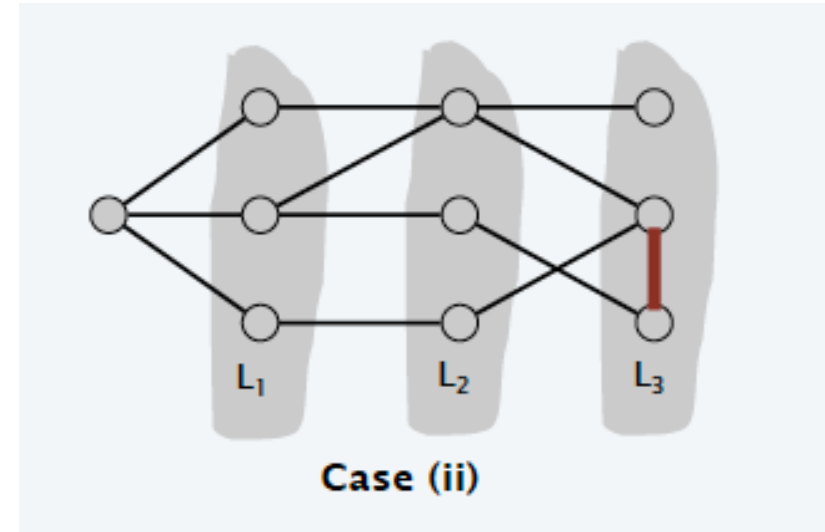
- Suppose no edge joins two nodes in same layer.
- By BFS property, *each edge joins two nodes in adjacent levels*.
- **Bipartition**: white = nodes on odd levels, blue = nodes on even levels.



BIPARTITE GRAPHS

Case ii:

- Suppose (x, y) is an edge with x, y in same level L_j .
- Let $z = \text{lca}(x, y) =$ lowest common ancestor.
- Let L_i be level containing z .
- Consider cycle that takes edge from x to y , then path from y to z , then path from z to x .
- Its length is $1 + (j - i) + (j - i) = 2(j - i) + 1$, which is odd. ■

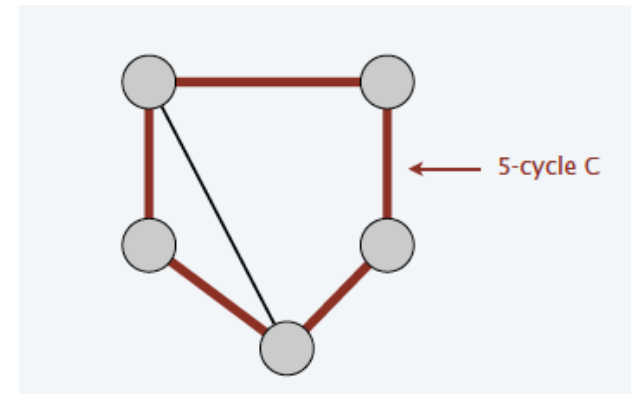
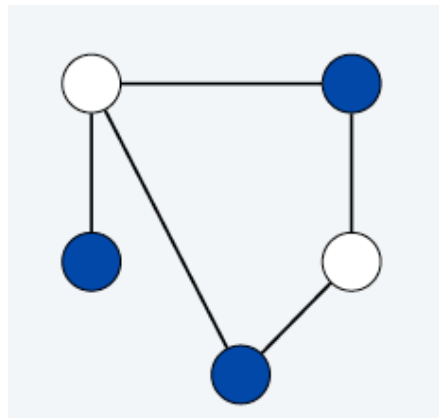


BIPARTITE GRAPHS

Lemma. Exactly one of the following holds.

- i. No edge of G joins two nodes of the same layer, and G is bipartite.
- ii. An edge of G joins two nodes of the same layer, and G contains an odd-length cycle (and hence is not bipartite).

Corollary. A graph G is bipartite if and only if it contains no odd-length cycle.



DIRECTED GRAPHS

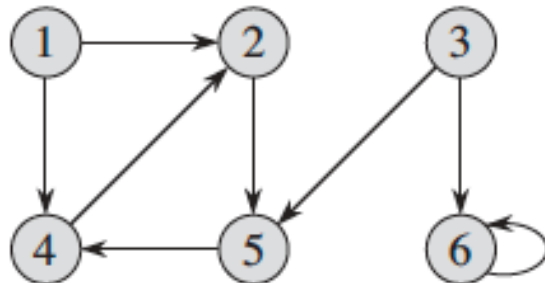
DIRECTED GRAPHS

Directed Graph $G = (V, E)$.

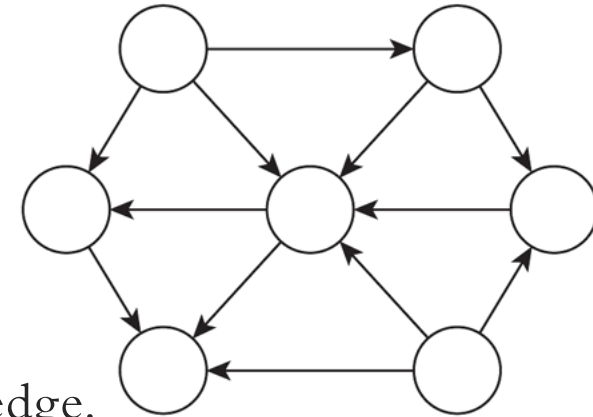
➤ Edge (u, v) leaves node u and **enters** node v .

Adjacency matrix. n -by- n matrix with $A_{uv} = 1$ if (u, v) is an edge.

- Space complexity $O(n^2)$.
- Checking if (u, v) is an edge takes constant time.
- Identifying all edges takes $\Theta(n^2)$ time.



In-degree: #incoming edges
Out-degree: #outgoing edges

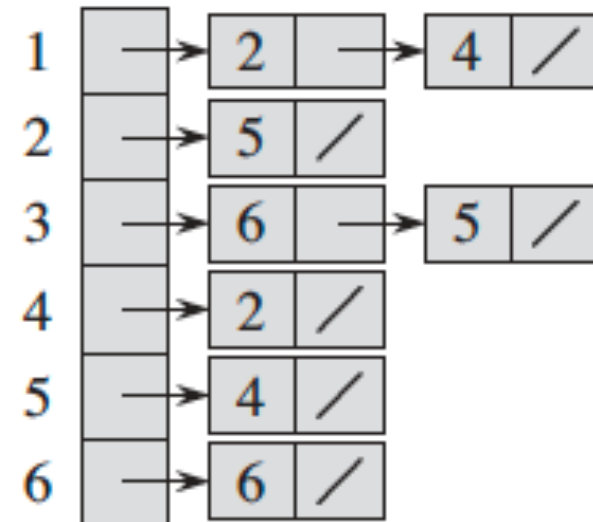
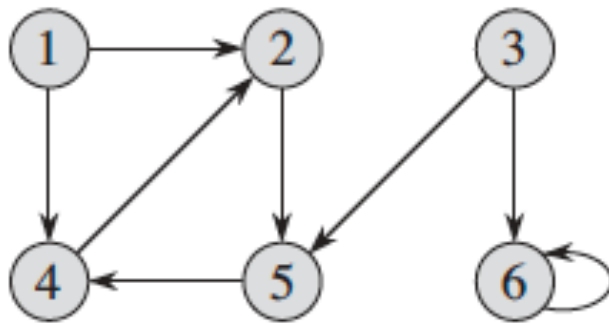


	1	2	3	4	5	6
1	0	1	0	1	0	0
2	0	0	0	0	1	0
3	0	0	0	0	1	1
4	0	1	0	0	0	0
5	0	0	0	1	0	0
6	0	0	0	0	0	1

DIRECTED GRAPHS

Adjacency lists. Node-indexed array of lists.

- Space complexity is $\Theta(m + n)$.
- Checking if (u, v) is an edge takes $O(\text{degree}(u))$ time.
- Identifying all edges takes $\Theta(m + n)$ time.



DIRECTED GRAPHS

Directed Graph $G = (V, E)$.

Directed reachability.

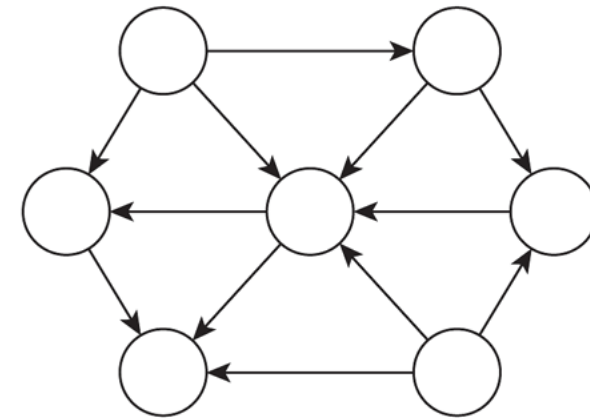
Given a node s , find all nodes reachable from s .

Directed $s - t$ shortest path problem.

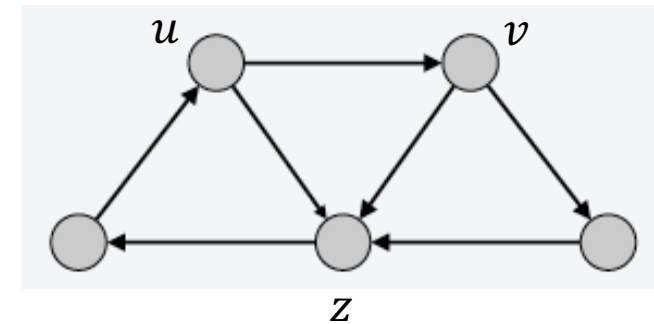
Given two nodes s and t , what is the length of a shortest path from s to t ?

➤ **Graph search:** BFS extends naturally to directed graphs.

Connectivity?



DIRECTED GRAPHS



Definition.

- Nodes u and v are **mutually reachable** if there is both a path from u to v and also a path from v to u .
- A graph is **strongly connected** if every pair of nodes is mutually reachable.

Lemma. Let s be any node. G is strongly connected if and only if **every node is reachable from s** , and **s is reachable from every node**.

Proof.

⇒ Follows from

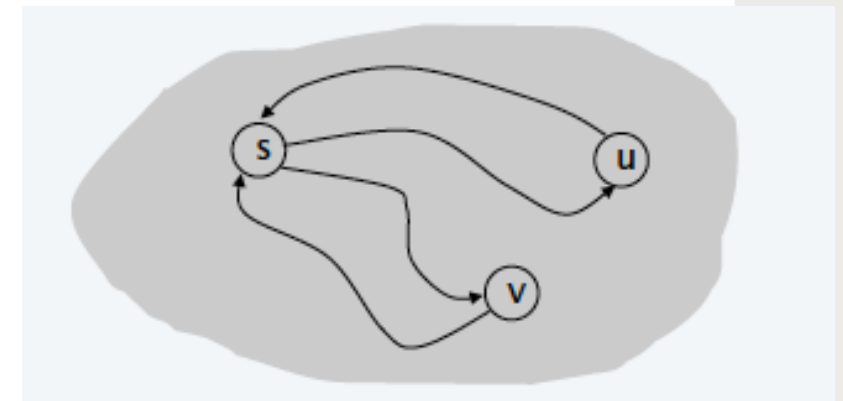
⇐ Path from s to

Path from s to

Given a graph, how to
determine whether it is
strongly connected or not?

with $s \rightsquigarrow v$ path.

with $s \rightsquigarrow u$ path.



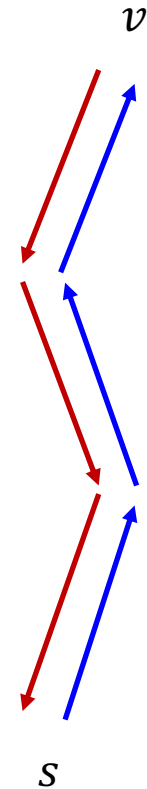
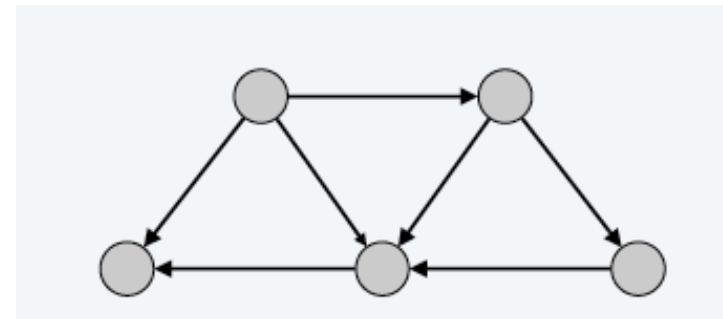
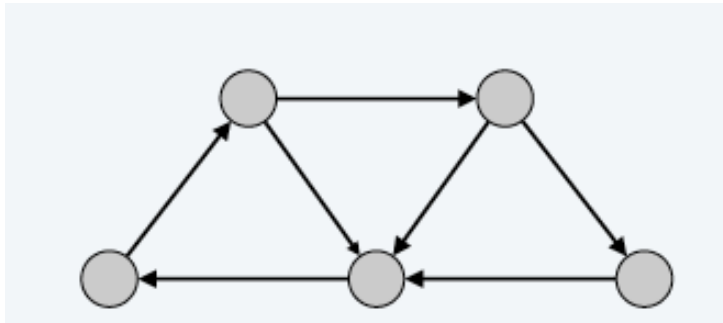
DIRECTED GRAPHS

Theorem. Can determine if G is strongly connected in $O(m + n)$ time.

Proof.

- Pick any node s .
- Run BFS from s in G .
- Run BFS from s in G^{reverse} .
- Return true if and only if all nodes reached in both BFS executions.
- Correctness follows immediately from previous lemma. ■

reverse orientation of every edge in G



DIRECTED ACYCLIC GRAPHS

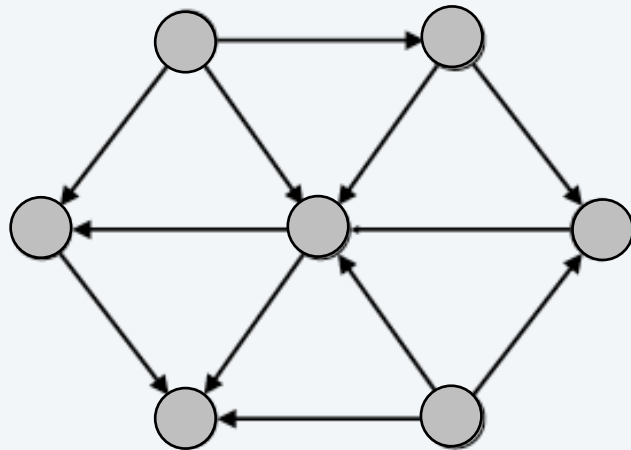
DIRECTED ACYCLIC GRAPHS

If the graph contains a cycle,
then no linear ordering is possible.

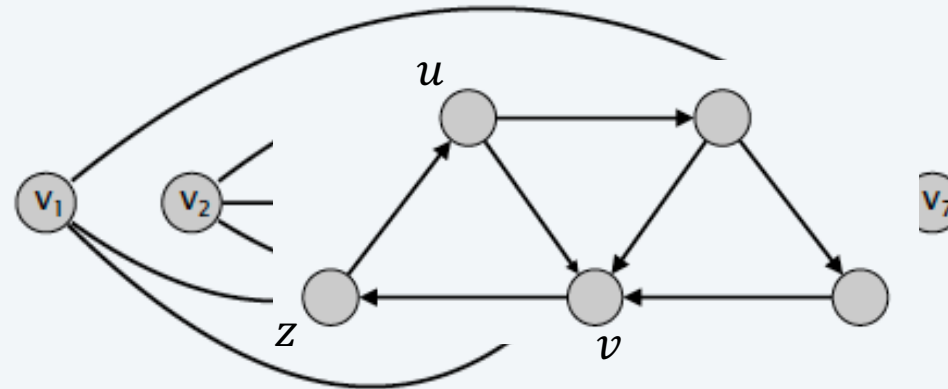
Definitions

An ordering of the nodes so that all edges point “forward”.

- A **directed acyclic graphs (DAG)** is a directed graph that contains no directed cycles.
- A **topological order** of a directed graph $G = (V, E)$ is an ordering of its nodes as v_1, v_2, \dots, v_n so that for every edge (v_i, v_j) we have $i < j$.



a DAG



a topological ordering

DIRECTED ACYCLIC GRAPHS



- How many topological orderings does the following graph have?
- An ordering of the nodes so that all edges point “forward”.

“a” must be first and “e” must be last

Exhaustive Search: there will be $5 \star 4 \star 3 \star 2 = 120$ possibilities.

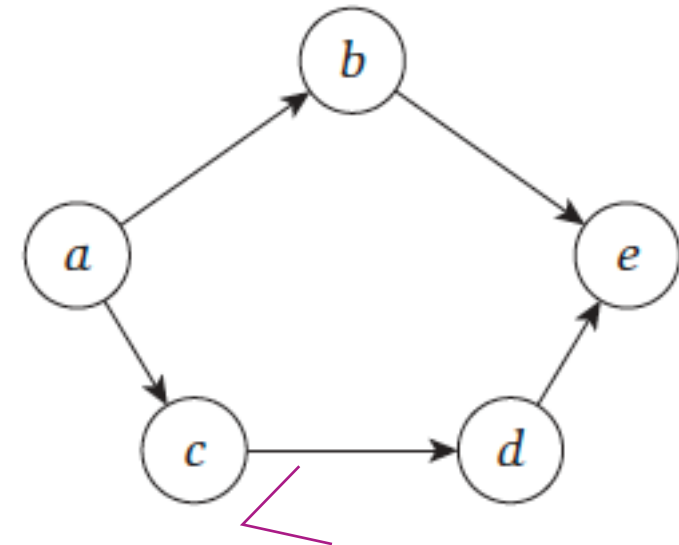
Observations:

- The first node must be one that has no edge coming into it.
- The last node must be one that has no edge leaving it.

a, **b**, c, d, e

a, c, **b**, d, e

a, c, d, **b**, e



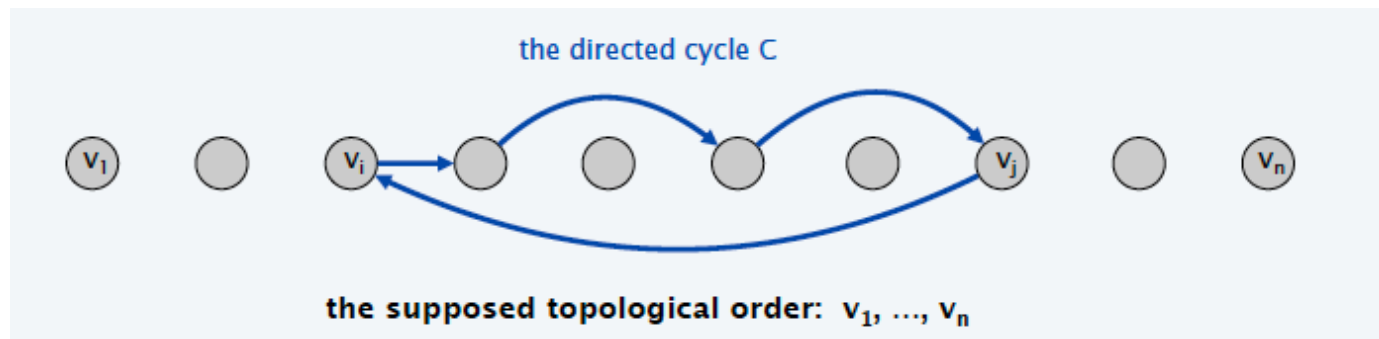
“c” must be before “d”

DIRECTED ACYCLIC GRAPHS

Lemma. If G has a topological order, then G is a DAG.

Proof [by contradiction]

- Suppose that G has a topological order v_1, v_2, \dots, v_n and that G also has a directed cycle C .
- Let v_i be the lowest-indexed node in C , and let v_j be the node just before v_i ;
thus (v_j, v_i) is an edge.
- By our choice of i , we have $i < j$.
- On the other hand, since (v_j, v_i) is an edge and v_1, v_2, \dots, v_n is a topological order, we must have $j < i$, a contradiction. ■



DIRECTED ACYCLIC GRAPHS

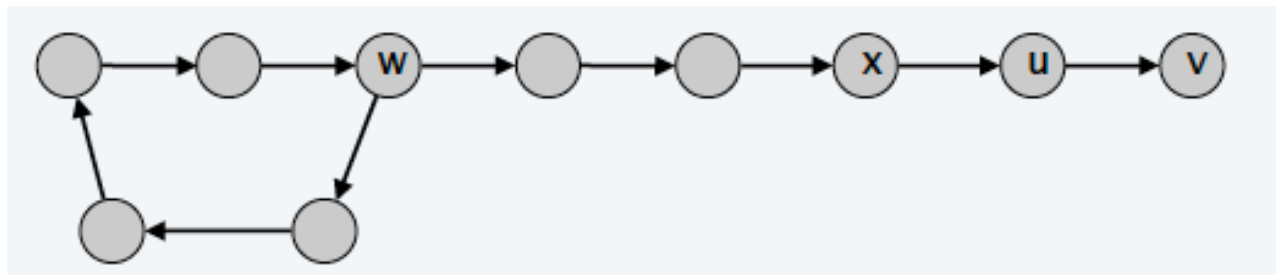
- Question: Does every DAG have a topological ordering?
- Question: If so, how do we compute one?

DIRECTED ACYCLIC GRAPHS

Lemma. If G is a DAG, then G has a node with no entering edges.

Proof.[by contradiction]

- Suppose that G is a DAG and every node has at least one entering edge.
- Pick any node v , and begin following edges backward from v .
- Since v has at least one entering edge (u, v) we can walk backward to u .
- Then, since u has at least one entering edge (x, u) , we can walk backward to x .
- Repeat until we visit a node, say w , twice.
- Let C denote the sequence of nodes encountered between successive visits to w .

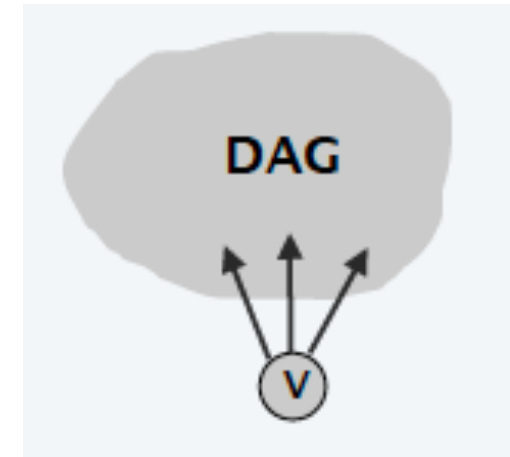


DIRECTED ACYCLIC GRAPHS

Lemma. If G is a DAG, then G has a topological ordering.

Proof. [by induction on n]

- Base case: true if $n = 1$.
- Inductive hypothesis: Assume true for $k < n$ nodes.
- Given DAG on $n > 1$ nodes, **find a node v with no entering edges.**
- $G - \{v\}$ is a DAG, since deleting v cannot create cycles.
- By inductive hypothesis, $G - \{v\}$ has a topological ordering.
- Place v first in topological ordering; then append nodes of $G - \{v\}$ in topological order.
- This is valid since v has no entering edges. ■



DIRECTED ACYCLIC GRAPHS

To compute a topological ordering of G :

Find a node v with no incoming edges and order it first

Delete v from G

Recursively compute a topological ordering of $G - \{v\}$
and append this order after v

Theorem. Algorithm finds a topological order in $O(m + n)$ time.

Proof:

Maintain the following information:

- $\text{count}(w)$ = remaining number of incoming edges to node w
- S = set of remaining nodes with no incoming edges

Initialization: $O(m + n)$ via single scan through graph.

DIRECTED ACYCLIC GRAPHS

To compute a topological ordering of G :

Find a node v with no incoming edges and order it first

Delete v from G

Recursively compute a topological ordering of $G - \{v\}$
and append this order after v

Update: to delete v

- remove v from S
- decrement $\text{count}(w)$ for all edges from v to w ; and add w to S if $\text{count}(w)$ hits 0
- this is $O(1)$ per edge ■

GREEDY ALGORITHMS

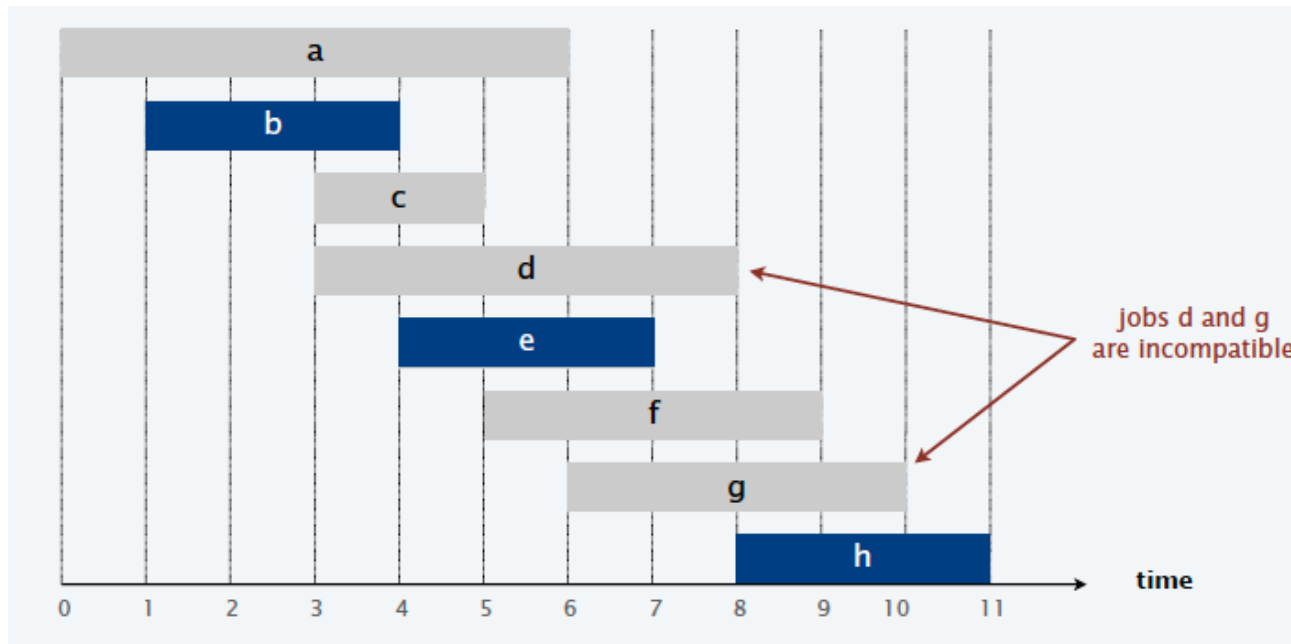
INTERVAL SCHEDULING PROBLEM

INTERVAL SCHEDULING PROBLEM

Given a set of jobs $J = \{1, 2, \dots, n\}$

- Job j starts at s_j and finishes at $f_j \geq s_j$.
- Two jobs (open intervals) are compatible if they don't overlap.

Goal: find maximum subset of mutually compatible jobs.



Intuition: shorter is better

INTERVAL SCHEDULING PROBLEM

Idea 1:

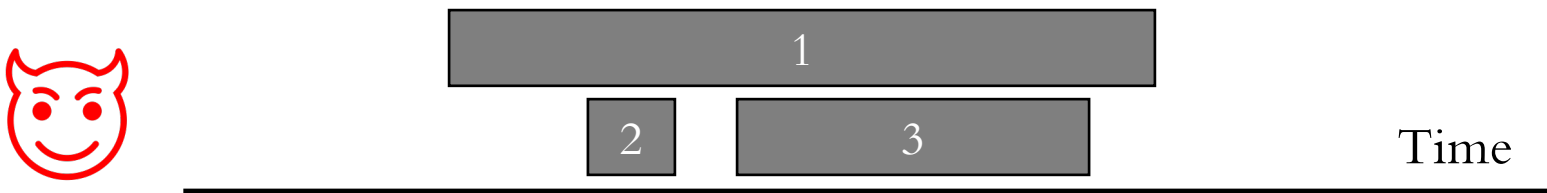
- Repeatedly pick **shortest** compatible, unscheduled job (i.e. that does not conflict with any scheduled job).



Idea 2:

Intuition: earlier is better

- Repeatedly pick compatible job with **earliest starting time**.



GREEDY ALGORITHM

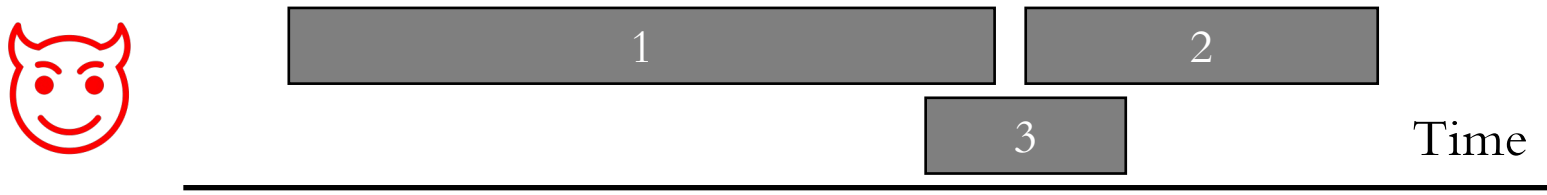
- Repeatedly pick an item until no more feasible choices.
- Among all feasible choices, we always pick the one that minimizes (or maximizes) some property.
 - length, starting time, ...
- Such algorithms are called *greedy*.
- Greedy algorithms may not be optimal.
- But maybe we have been using the wrong property!

INTERVAL SCHEDULING PROBLEM

What about earliest-finish-time-first?

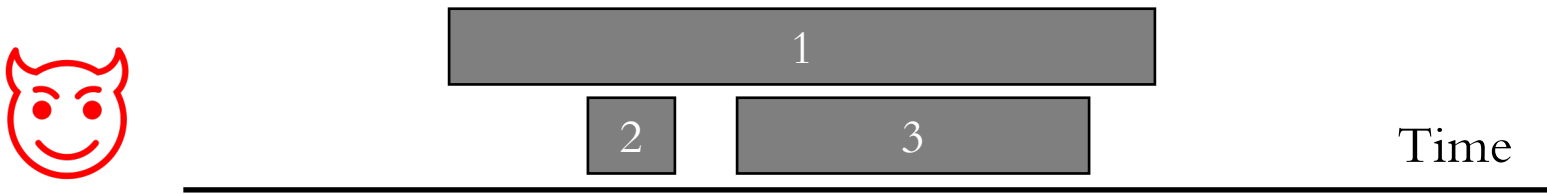
Idea 1:

- Repeatedly pick **shortest** compatible, unscheduled job (i.e. that does not conflict with any scheduled job).



Idea 2:

- Repeatedly pick compatible job with **earliest starting time**.



EARLIEST-FINISH-TIME-FIRST ALGORITHM

EARLIEST-FINISH-TIME-FIRST ($n, s_1, s_2, \dots, s_n, f_1, f_2, \dots, f_n$)

SORT jobs by finish times and renumber so that $f_1 \leq f_2 \leq \dots \leq f_n$.

$S \leftarrow \emptyset$. \leftarrow set of jobs selected

FOR $j = 1$ **TO** n

IF (job j is compatible with S)

$S \leftarrow S \cup \{ j \}$.

RETURN S .

Proposition. Can implement earliest-finish-time first in $O(n \log n)$ time.

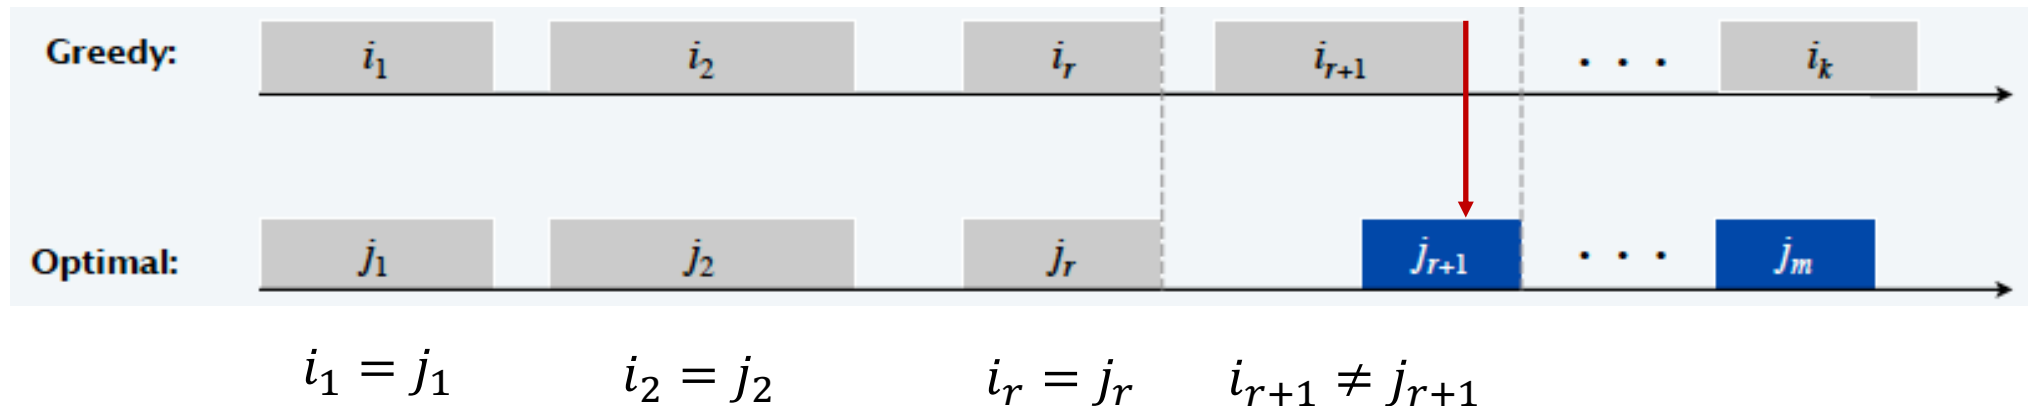
EARLIEST-FINISH-TIME-FIRST ALGORITHM

Theorem. The earliest-finish-time-first algorithm is optimal.

Proof. [by contradiction]

- Assume Greedy is not optimal.
- Let $A = \{i_1, i_2, \dots, i_k\}$ be set of jobs selected by Greedy.
- Let $O = \{j_1, j_2, \dots, j_m\}$ be set of jobs in an optimal solution. Then $m > k$.
- Let $r + 1$ be first index such that $i_{r+1} \neq j_{r+1}$. such a job exists $\Rightarrow f_{i_{r+1}} \leq f_{j_{r+1}}$

Switching j_{r+1} by i_{r+1} in O :
Still *feasible* and *optimal*!



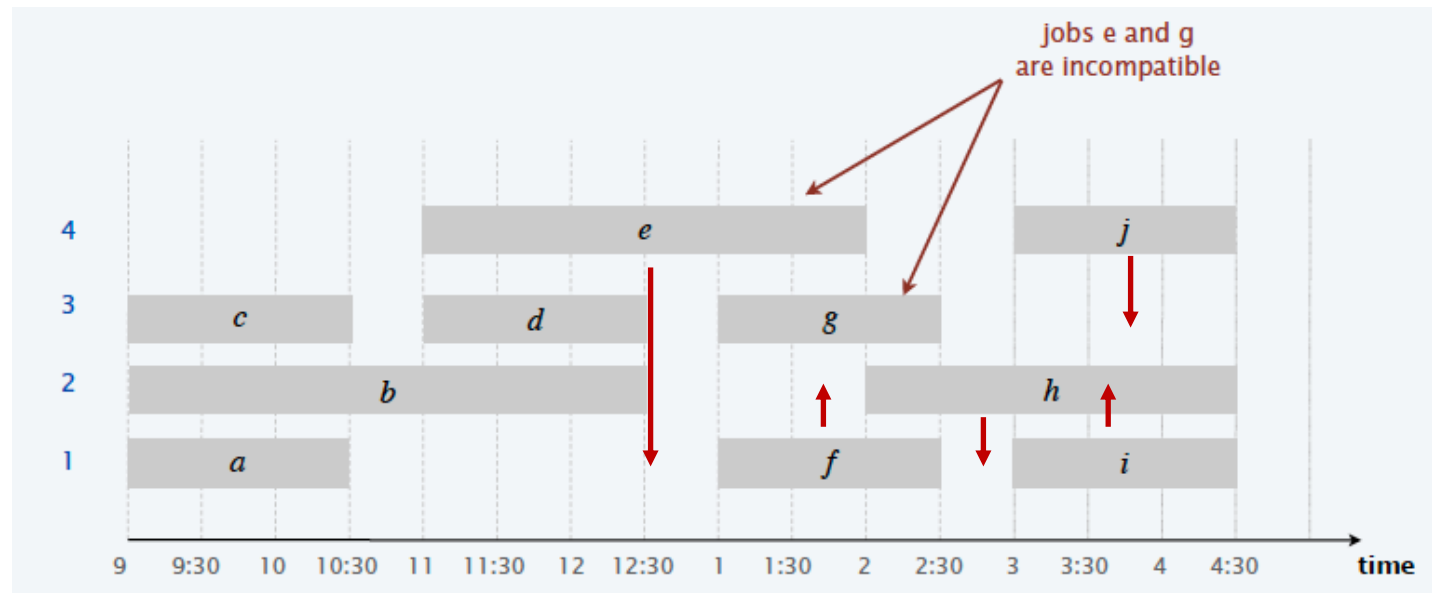
INTERVAL PARTITIONING

INTERVAL PARTITIONING

Given a set of lectures (jobs) $L = \{1, 2, \dots, n\}$;

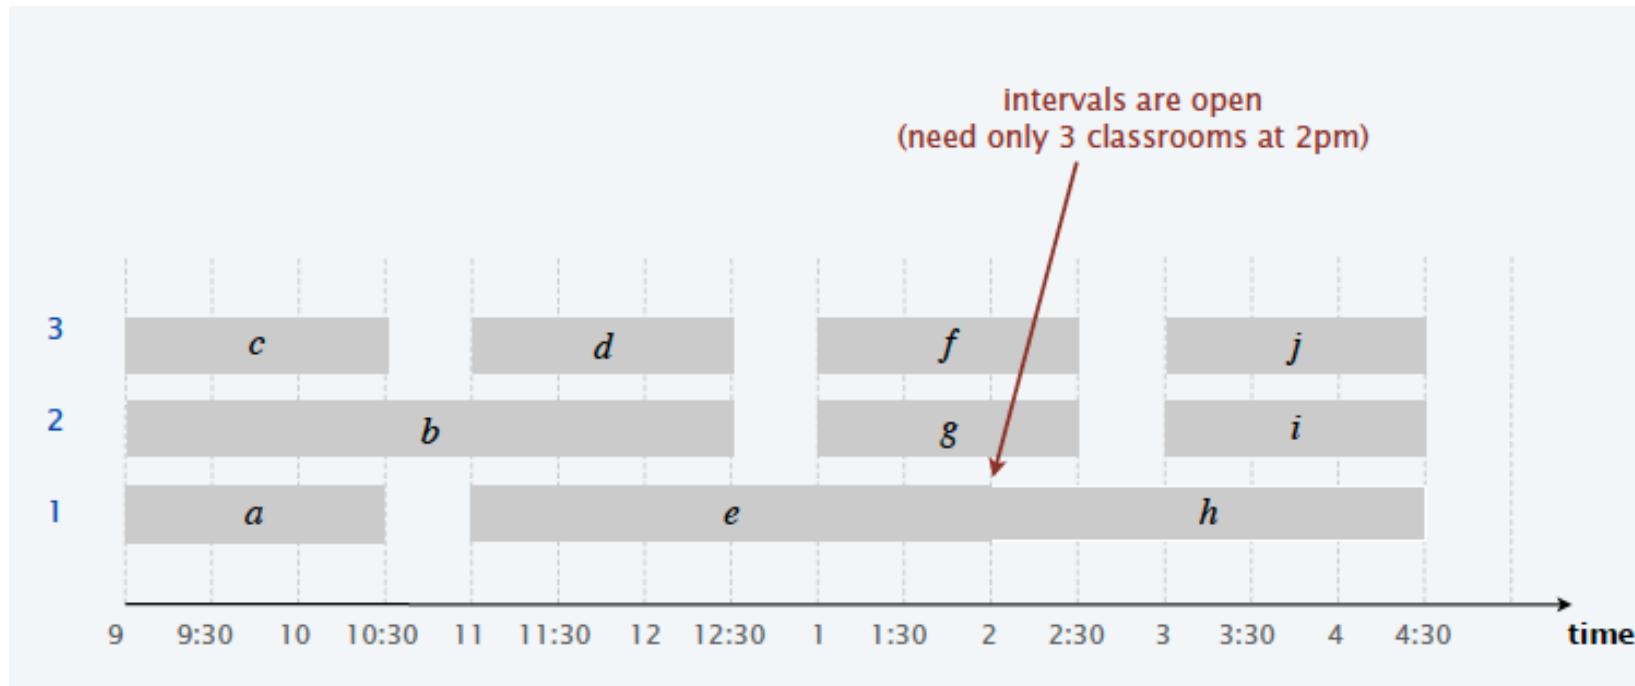
- Lecture j starts at s_j and finishes at $f_j \geq s_j$.
- Two lectures are compatible if they don't overlap.

Goal: find minimum number of classrooms to schedule all lectures so that no two lectures occur at the same time in the same room



INTERVAL PARTITIONING

- Optimal is 3 classrooms.

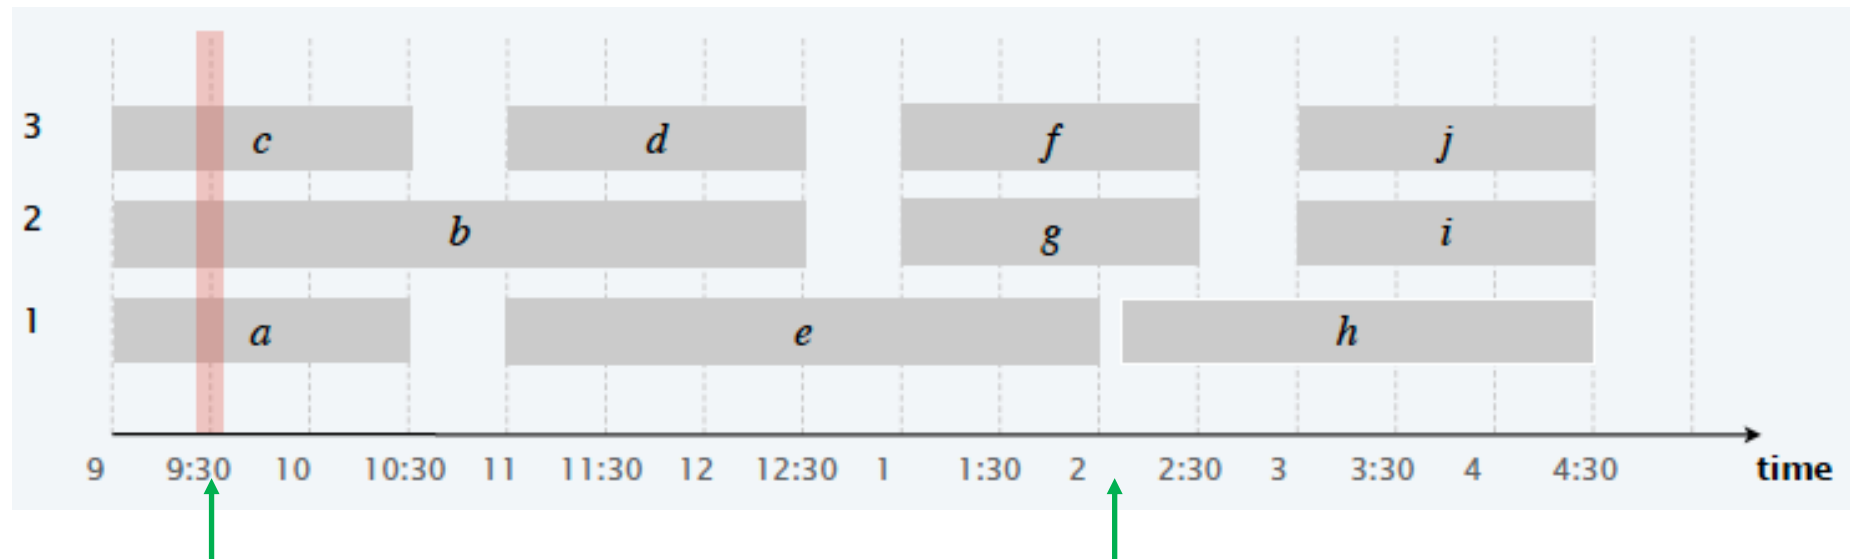


INTERVAL PARTITIONING

Definition. The depth of a set of open intervals is the maximum number of intervals that contain any given point.

Key observation. #rooms needed \geq depth.

Is depth enough???

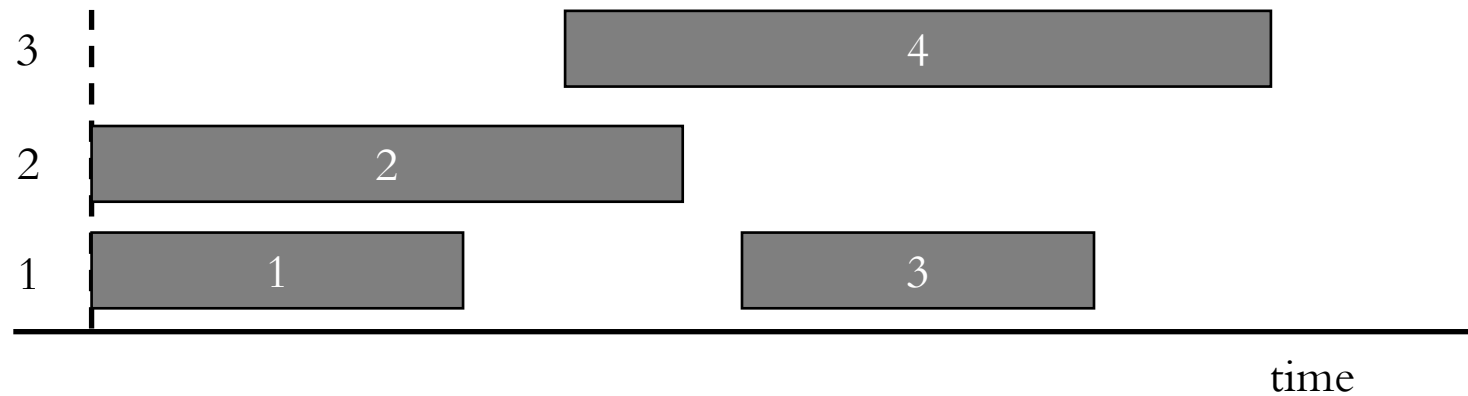
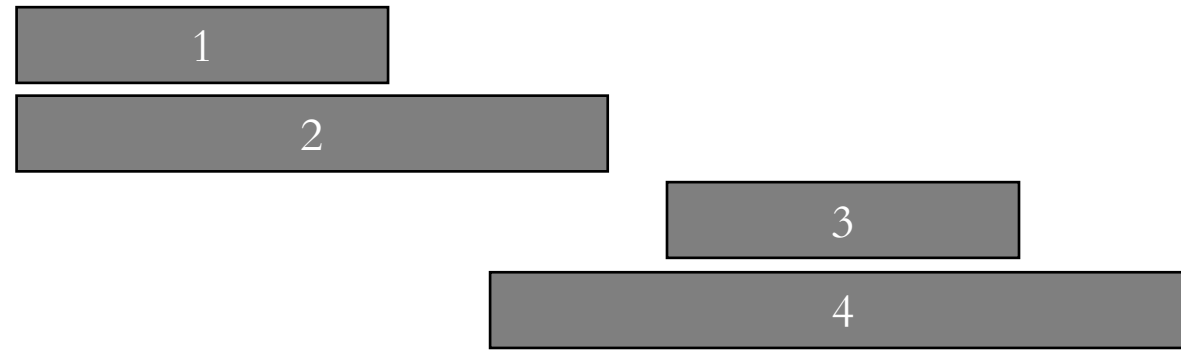


3 classrooms are needed

2 classrooms are needed

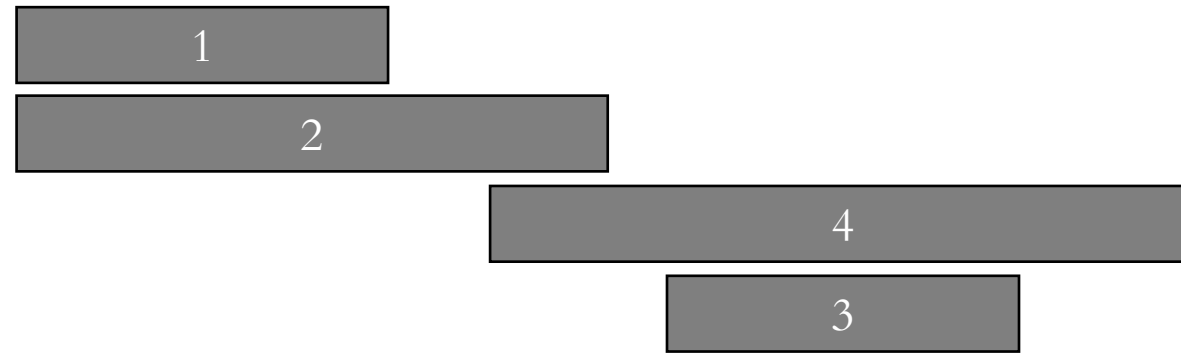
INTERVAL PARTITIONING

Can we do earliest-**finish**-time-first?

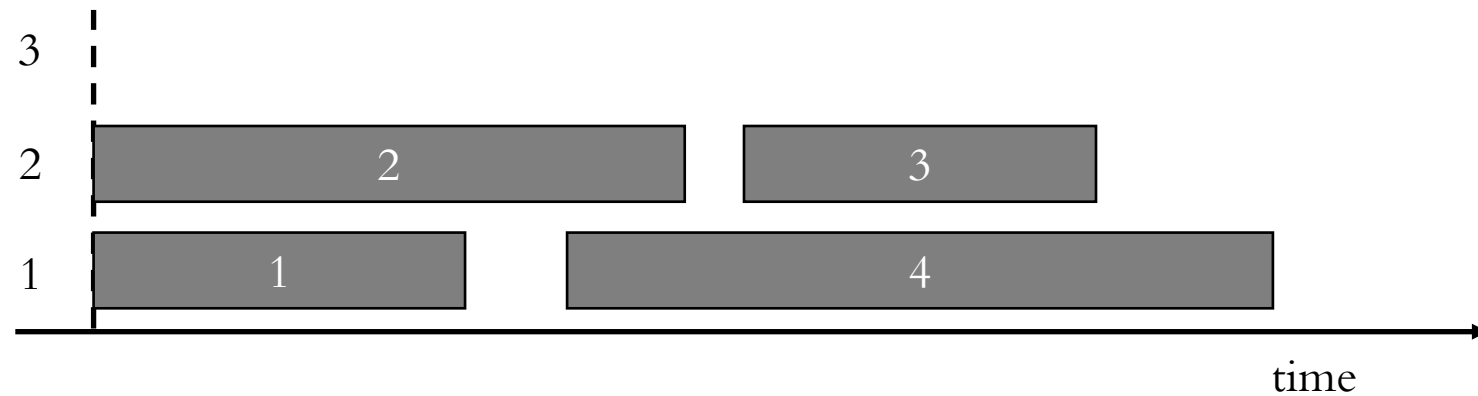


INTERVAL PARTITIONING

Can we do earliest-**start**-time-first?



40



INTERVAL PARTITIONING: EARLIEST-START-TIME-FIRST ALGORITHM

EARLIEST-START-TIME-FIRST ($n, s_1, s_2, \dots, s_n, f_1, f_2, \dots, f_n$)

SORT lectures by start times and renumber so that $s_1 \leq s_2 \leq \dots \leq s_n$.

$d \leftarrow 0$. \leftarrow number of allocated classrooms

FOR $j = 1$ TO n

IF (lecture j is compatible with some classroom)

Schedule lecture j in any such classroom k .

ELSE

Allocate a new classroom $d + 1$.

Schedule lecture j in classroom $d + 1$.

$d \leftarrow d + 1$.

RETURN schedule.

Lemma.

The earliest-start-time-first algorithm can be implemented in $O(n \log n)$ time.

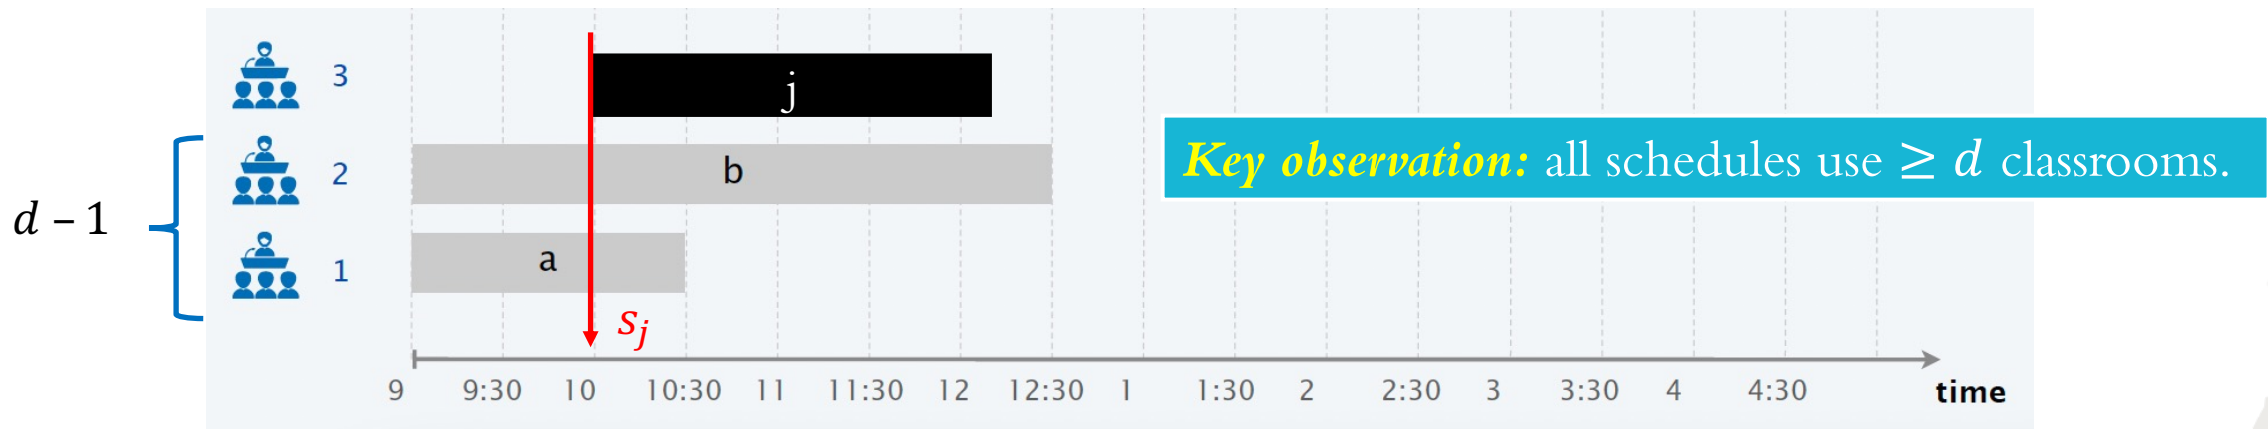
Lemma.

The earliest-start-time first algorithm never schedules two incompatible lectures in the same classroom.

INTERVAL PARTITIONING: EARLIEST-START-TIME-FIRST ALGORITHM

Theorem. Earliest-start-time-first algorithm uses #depth rooms and thus is optimal.

- Let d = number of classrooms that the algorithm allocates.
- Classroom d is opened because we needed to schedule a lecture, say j , that is incompatible with a lecture in each of $d - 1$ other classrooms.
- Thus, these d lectures each end after s_j . → The d lectures are incompatible.
- Since we sorted by start time, each of these incompatible lectures start no later than s_j . ■



SCHEDULING TO MINIMIZING LATENESS

SCHEDULING TO MINIMIZING LATENESS

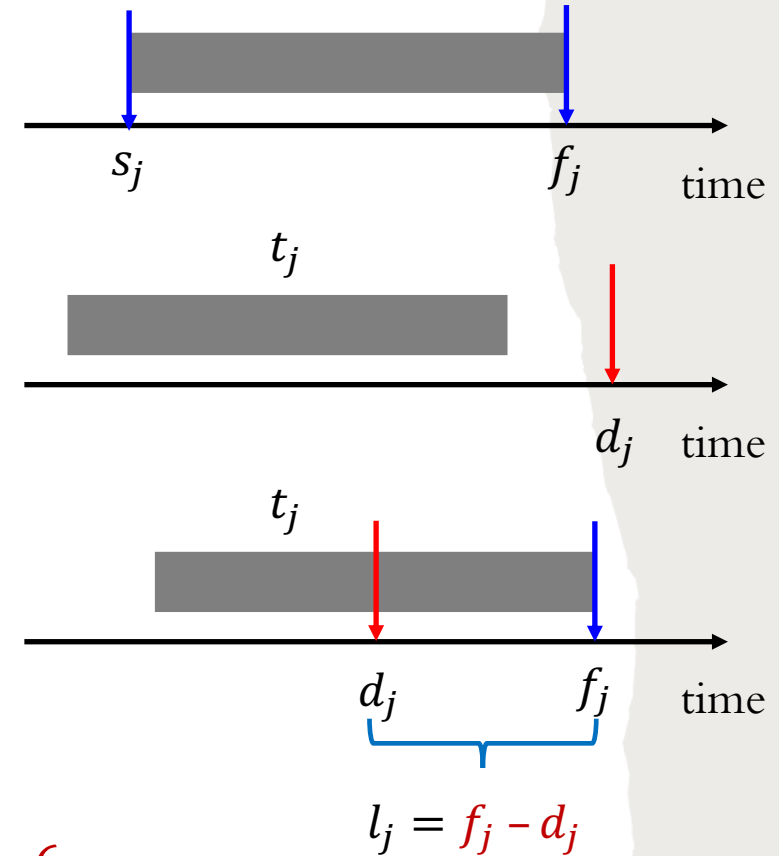
Single resource processes one job at a time.

➤ Job j requires t_j units of processing time and is due at time d_j .

➤ If j starts at time s_j , it finishes at time $f_j = s_j + t_j$.

➤ **Lateness:** $l_j = \max\{0, f_j - d_j\}$.

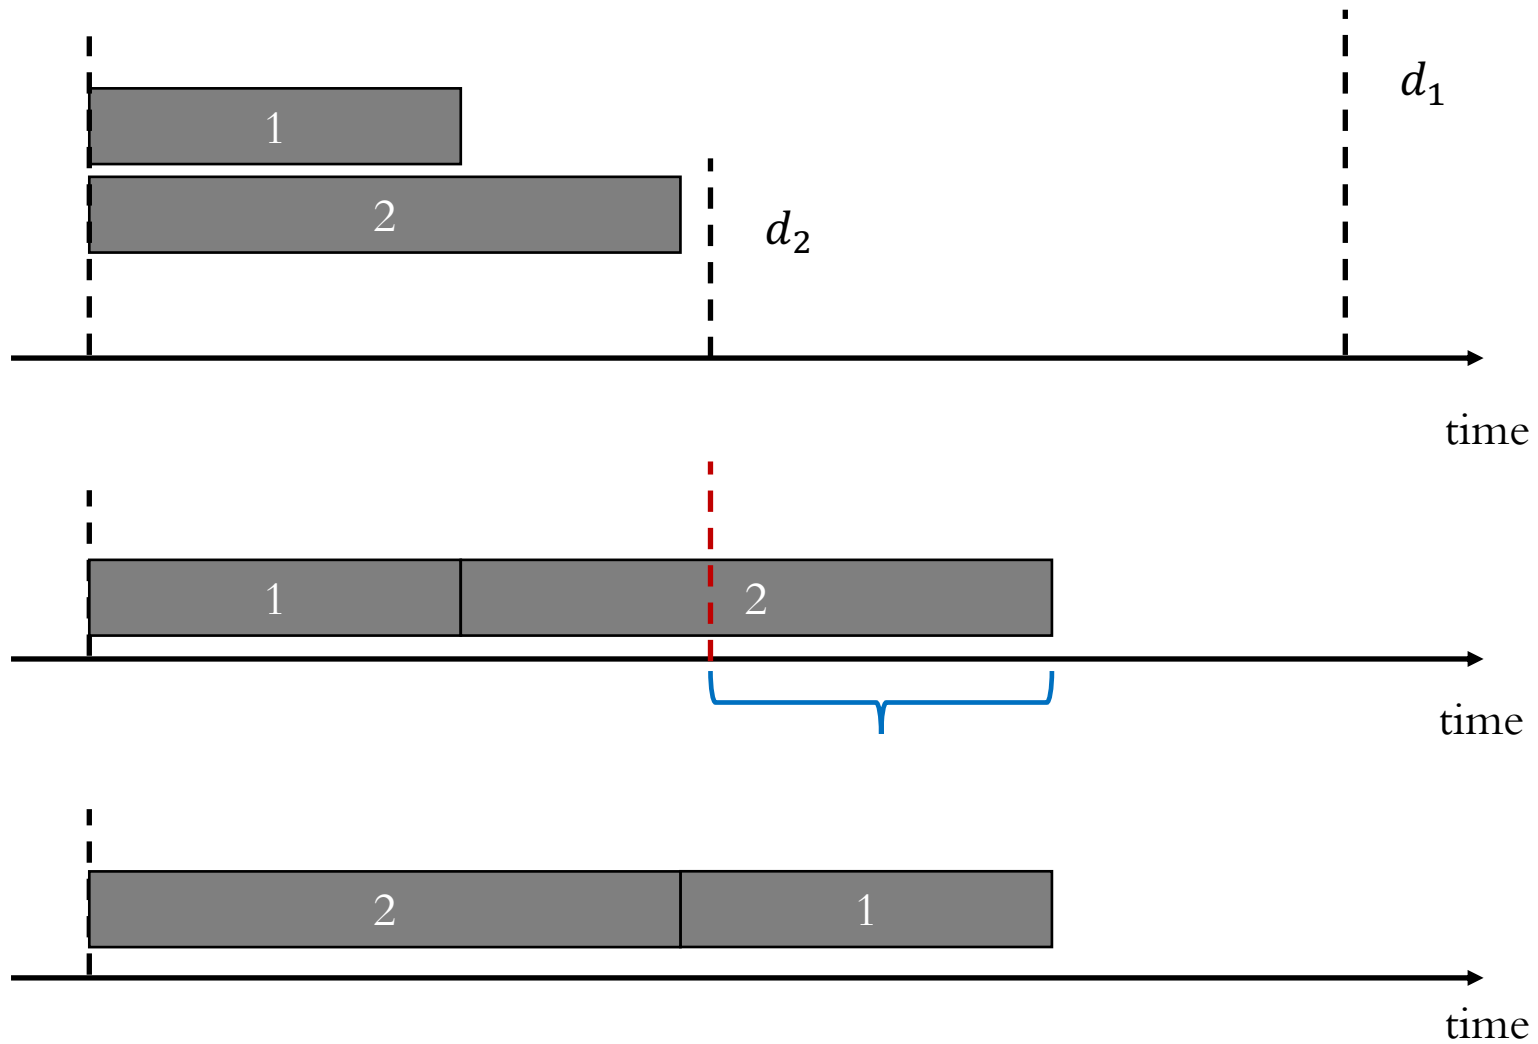
Goal: schedule all jobs to minimize **maximum** lateness $L = \max_j l_j$.



	1	2	3	4	5	6
t_j	3	2	1	4	3	2
d_j	6	8	9	9	14	15



SCHEDULING TO MINIMIZING LATENCY



SCHEDULING TO MINIMIZING LATENESS

EARLIEST-DEADLINE-FIRST ($n, t_1, t_2, \dots, t_n, d_1, d_2, \dots, d_n$)

SORT jobs by due times and renumber so that $d_1 \leq d_2 \leq \dots \leq d_n$.

$t \leftarrow 0$.

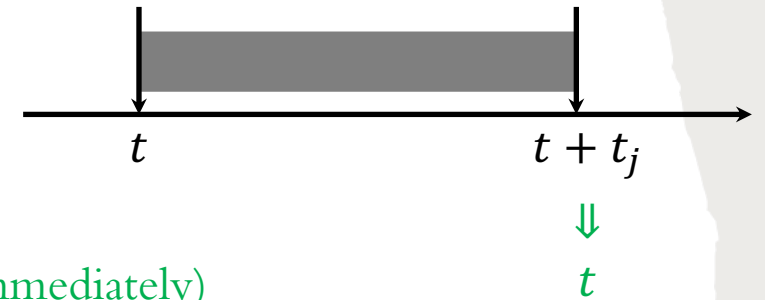
FOR $j = 1$ **TO** n ← Process the ordered jobs one by one (immediately)

Assign job j to interval $[t, t + t_j]$.

$s_j \leftarrow t$; $f_j \leftarrow t + t_j$.

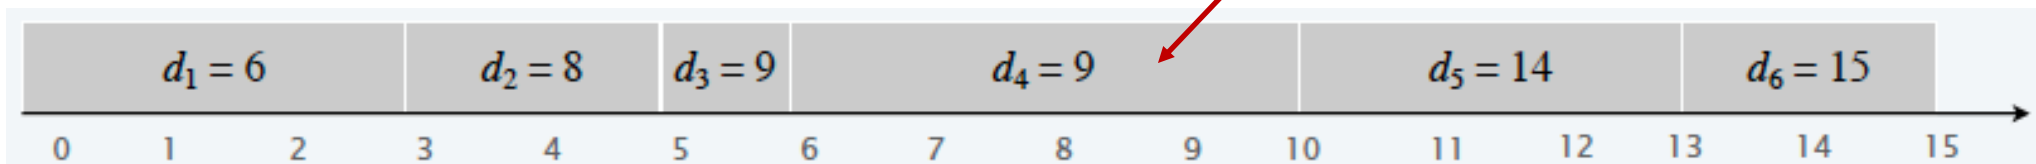
$t \leftarrow t + t_j$.

RETURN intervals $[s_1, f_1], [s_2, f_2], \dots, [s_n, f_n]$.



	1	2	3	4	5	6
t_j	3	2	1	4	3	2
d_j	6	8	9	9	14	15

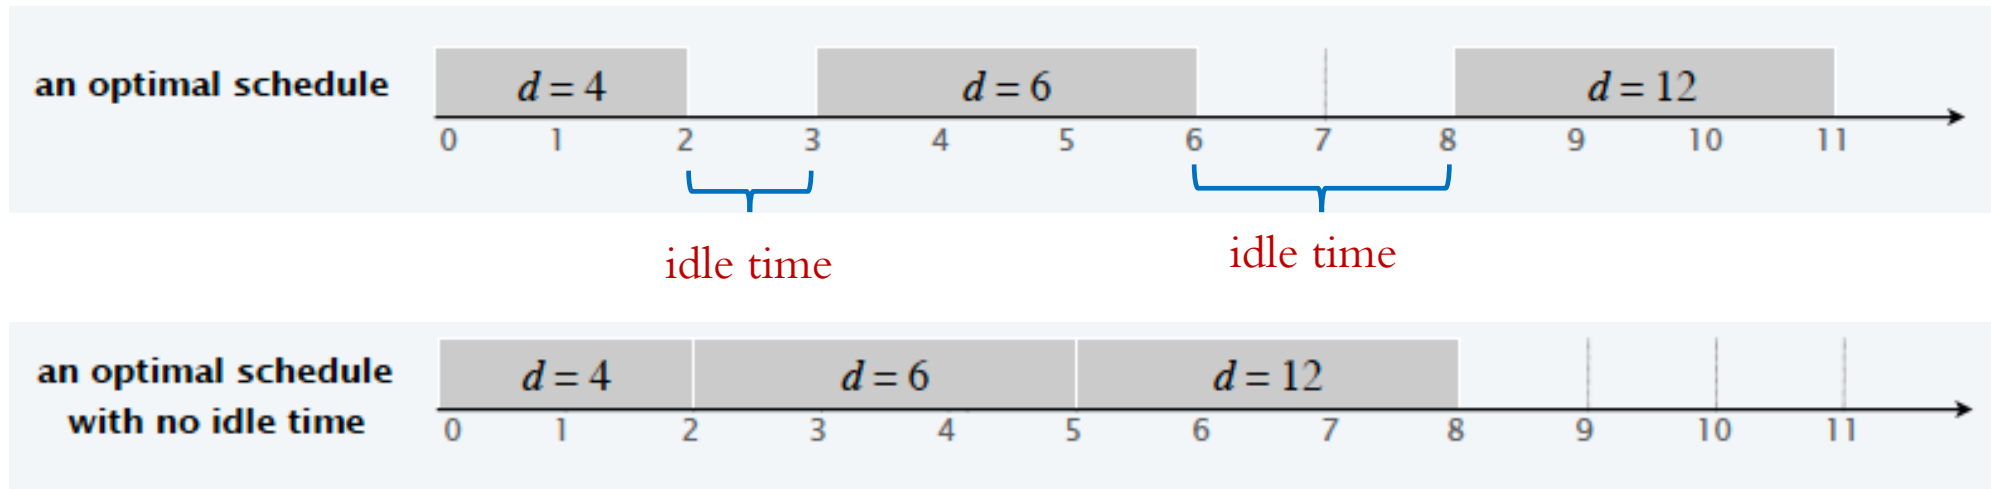
$l_4 = 1$



SCHEDULING TO MINIMIZING LATENESS

Properties for optimal schedules.

Observation 1. There exists an optimal schedule with no idle time.

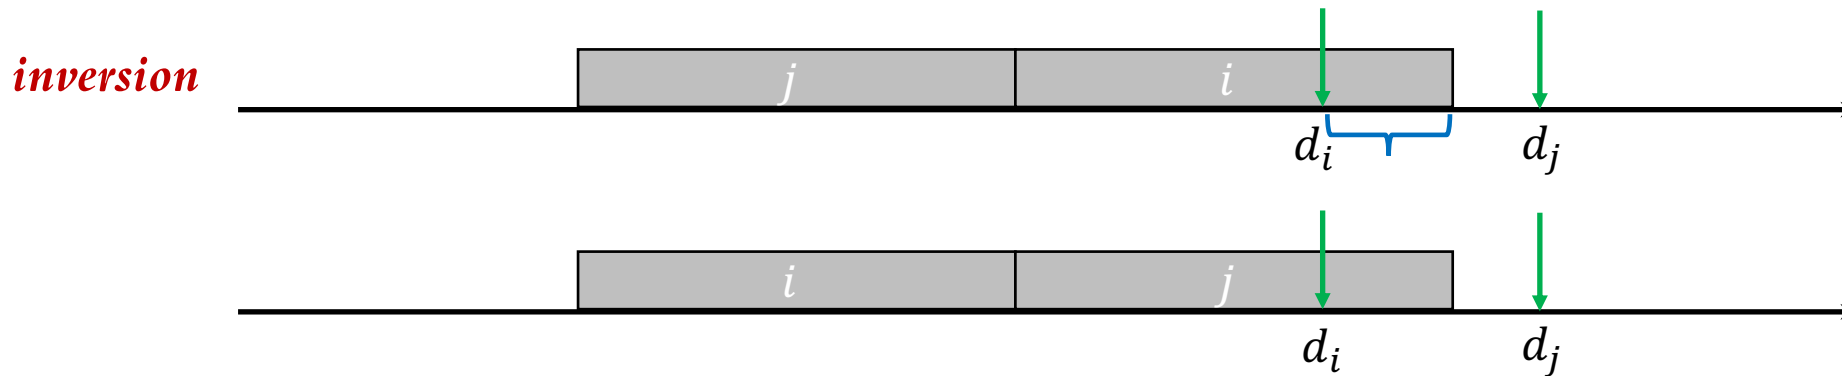


Observation 2. The earliest-deadline-first schedule has no idle time.

SCHEDULING TO MINIMIZING LATENESS

Definition. Given a schedule S , an **inversion** is a pair of jobs i and j such that: $d_i < d_j$ but j is scheduled before i .

or $i < j$ for ordered jobs



swap makes the schedule better!

Observation 3. The earliest-deadline-first schedule is the **unique** idle-free schedule with no inversions.

SCHEDULING TO MINIMIZING LATENESS

Observation 4. If an idle-free schedule has an inversion, then it has an **adjacent inversion**.

two inverted jobs scheduled consecutively

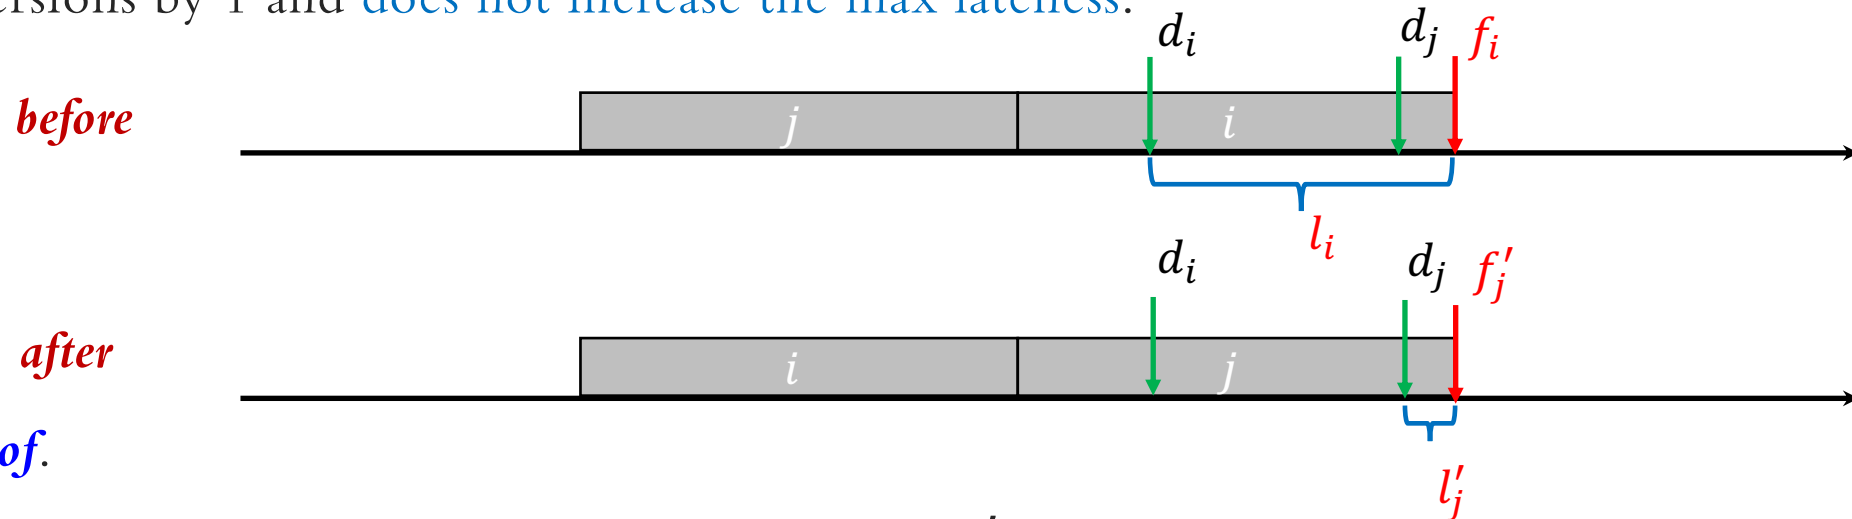
Proof.

- Let $i-j$ be a **closest** inversion. $d_j > d_i$
- Let k be element immediately to the right of j .
 - **Case 1:** $d_j > d_k$. Then $j-k$ is an adjacent inversion.
 - **Case 2.** $d_j < d_k$. Then $i-k$ is a closer inversion. ■



SCHEDULING TO MINIMIZING LATENESS

Key Claim. Exchanging two **adjacent**, **inverted** jobs i and j reduces the number of inversions by 1 and **does not increase the max lateness**.



Proof.

- Let l be the lateness before the swap, and let l' be it afterwards.
- $l'_k = l_k$ for all $k \neq i, j$.
- $l'_i \leq l_i$
- If job j is late, $l'_j = f'_j - d_j = f_i - d_j \leq f_i - d_i \leq l_i$. ■

SCHEDULING TO MINIMIZING LATENESS

Theorem. The earliest-deadline-first schedule S is optimal.

Proof. [by contradiction]

- Define S^* to be an optimal schedule with the **fewest inversions**.
- Can assume S^* has no idle time. → Observation 1
- **Case 1:** S^* has no inversions. Then $S = S^*$. → Observation 3
- **Case 2:** S^* has an inversion.
 - Let $i - j$ be an **adjacent** inversion → Observation 4
 - Exchanging jobs i and j decreases the number of inversions by 1 without increasing the max lateness → Key Claim
 - Contradicts “**fewest inversions**” part of the definition of S^* . ■

GREEDY ANALYSIS STRATEGIES

Greedy algorithm stays ahead.

- Show that after each step of the greedy algorithm, its solution is **at least as good as any other algorithm's**.
- [Interval scheduling]

Structural.

- Discover a simple “structural” bound asserting that **every possible solution must have a certain value**. Then show that your algorithm always achieves this bound.
- [Interval partitioning]

Exchange argument.

- Gradually **transform any solution to the one found by the greedy algorithm** without hurting its quality.
- [Minimizing lateness, Interval scheduling]

Thank You!