

COMP 3011
DESIGN AND ANALYSIS OF ALGORITHMS
FALL 2024

Greedy Algorithms

LI Bo
Department of Computing
The Hong Kong Polytechnic University

GREEDY ALGORITHMS

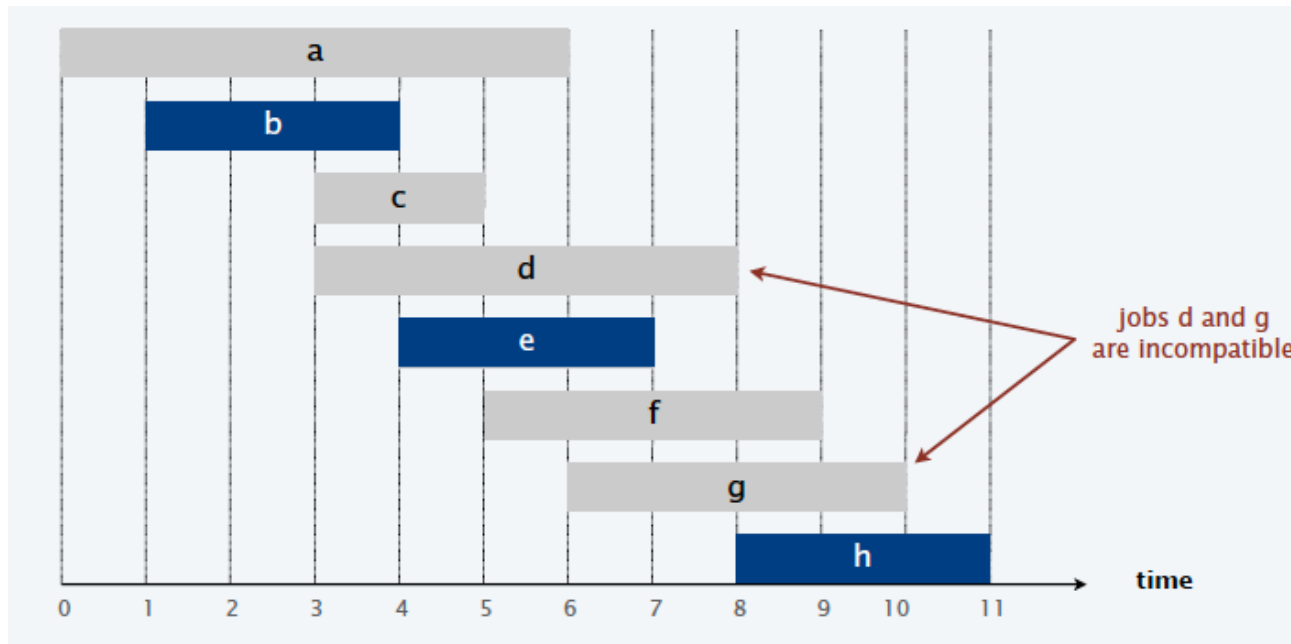
INTERVAL SCHEDULING PROBLEM

INTERVAL SCHEDULING PROBLEM

Given a set of jobs $J = \{1, 2, \dots, n\}$

- Job j starts at s_j and finishes at $f_j \geq s_j$.
- Two jobs (open intervals) are compatible if they don't overlap.

Goal: find maximum subset of mutually compatible jobs.



Intuition: shorter is better

INTERVAL SCHEDULING PROBLEM

Idea 1:


- Repeatedly pick **shortest** compatible, unscheduled job (i.e. that does not conflict with any scheduled job).



Idea 2:

Intuition: earlier is better

- Repeatedly pick compatible job with **earliest starting time**.



GREEDY ALGORITHM

- Repeatedly pick an item until no more feasible choices.
- Among all feasible choices, we always pick the one that minimizes (or maximizes) some property.
 - length, starting time, ...
- Such algorithms are called *greedy*.
- Greedy algorithms may not be optimal.
- But maybe we have been using the wrong property!

INTERVAL SCHEDULING PROBLEM

What about earliest-finish-time-first?

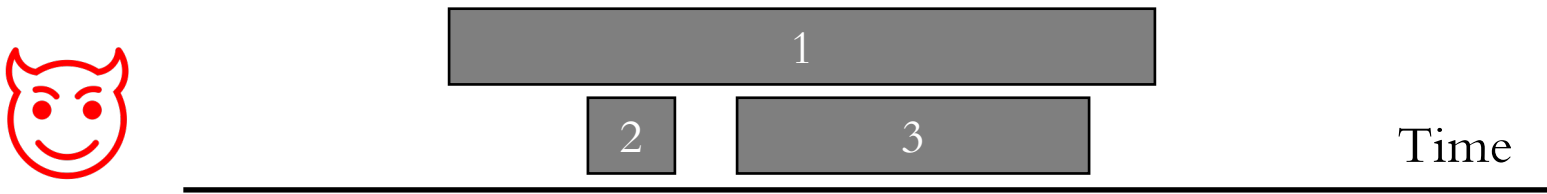
Idea 1:

- Repeatedly pick **shortest** compatible, unscheduled job (i.e. that does not conflict with any scheduled job).



Idea 2:

- Repeatedly pick compatible job with **earliest starting time**.



EARLIEST-FINISH-TIME-FIRST ALGORITHM

EARLIEST-FINISH-TIME-FIRST ($n, s_1, s_2, \dots, s_n, f_1, f_2, \dots, f_n$)

SORT jobs by finish times and renumber so that $f_1 \leq f_2 \leq \dots \leq f_n$.

$S \leftarrow \emptyset$. \leftarrow set of jobs selected

FOR $j = 1$ **TO** n

IF (job j is compatible with S)

$S \leftarrow S \cup \{ j \}$.

RETURN S .

Proposition. Can implement earliest-finish-time first in $O(n \log n)$ time.

EARLIEST-FINISH-TIME-FIRST ALGORITHM


Theorem. The earliest-finish-time-first algorithm is optimal.

Proof. [by contradiction]

- Assume Greedy is not optimal.
- Let $A = \{i_1, i_2, \dots, i_k\}$ be set of jobs selected by Greedy.
- Let $O = \{j_1, j_2, \dots, j_m\}$ be set of jobs in an optimal solution. Then $m > k$.
- Let $r + 1$ be first index such that $i_{r+1} \neq j_{r+1}$.

Switching j_{r+1} by i_{r+1} in O :
Still *feasible* and *optimal*!

such a job exists $\rightarrow f_{i_{r+1}} \leq f_{j_{r+1}}$



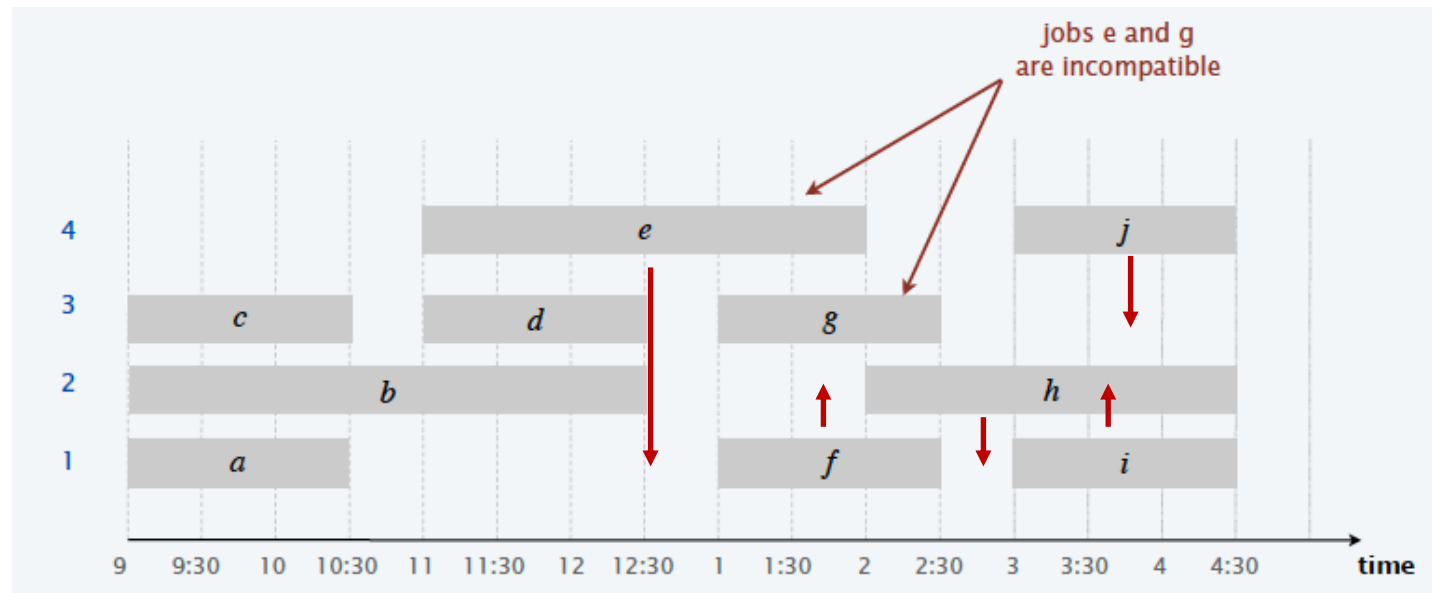
INTERVAL PARTITIONING

INTERVAL PARTITIONING

Given a set of lectures (jobs) $L = \{1, 2, \dots, n\}$;


- Lecture j starts at s_j and finishes at $f_j \geq s_j$.
- Two lectures are compatible if they don't overlap.

Goal: find minimum number of classrooms to schedule all lectures so that no two lectures occur at the same time in the same room



INTERVAL PARTITIONING

- Optimal is 3 classrooms.




INTERVAL PARTITIONING

Definition. The depth of a set of open intervals is the maximum number of intervals that contain any given point.

Key observation. #rooms needed \geq depth.

Is depth enough???


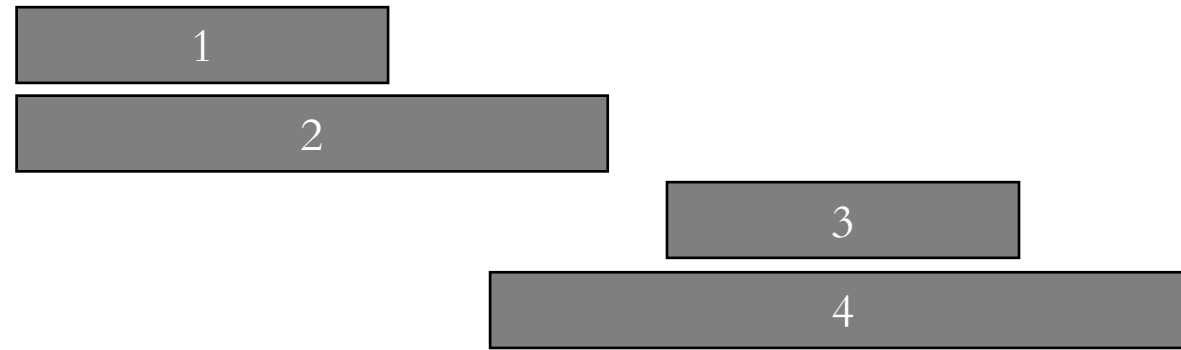


3 classrooms are needed

2 classrooms are needed

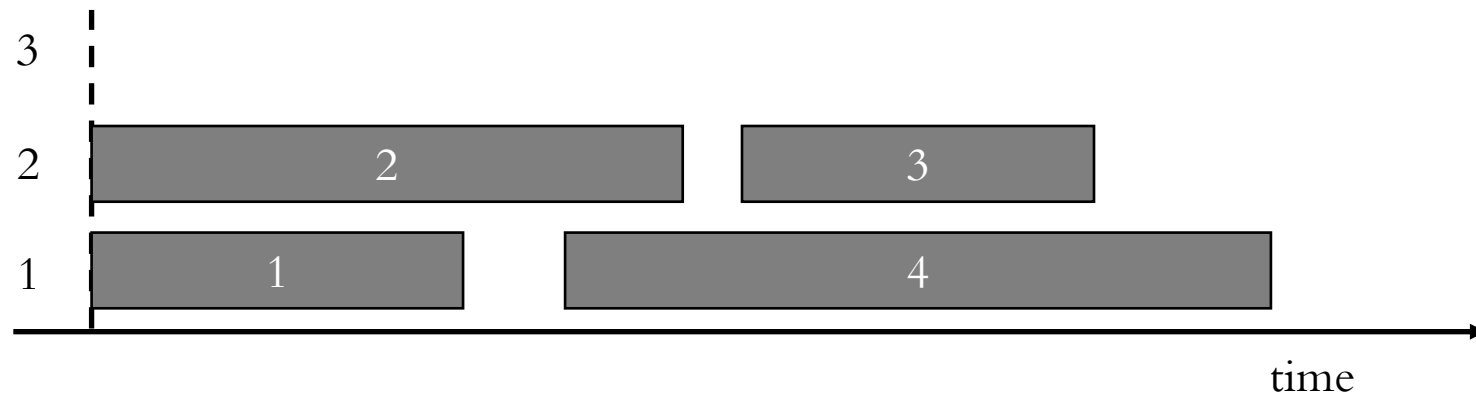
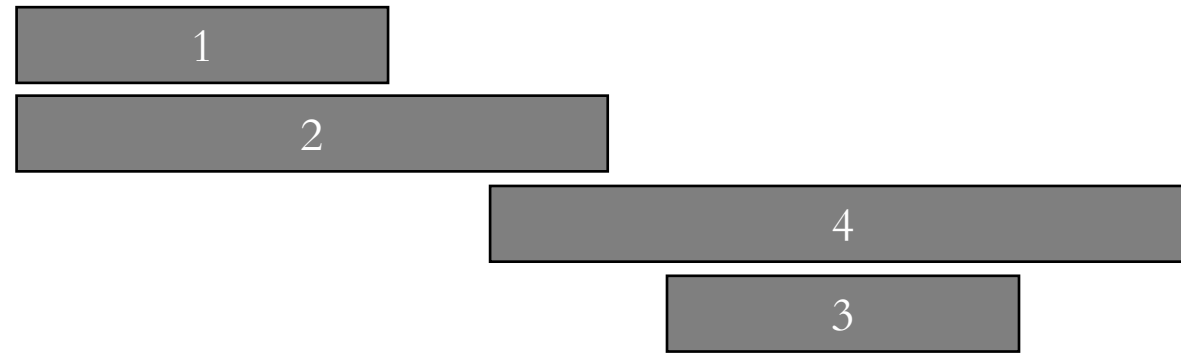
INTERVAL PARTITIONING

Can we do earliest-**finish**-time-first?



INTERVAL PARTITIONING

Can we do earliest-**start**-time-first?



INTERVAL PARTITIONING: EARLIEST-START-TIME-FIRST ALGORITHM

EARLIEST-START-TIME-FIRST ($n, s_1, s_2, \dots, s_n, f_1, f_2, \dots, f_n$)

SORT lectures by start times and renumber so that $s_1 \leq s_2 \leq \dots \leq s_n$.

$d \leftarrow 0$. \leftarrow number of allocated classrooms

FOR $j = 1$ TO n

IF (lecture j is compatible with some classroom)

Schedule lecture j in any such classroom k .

ELSE

Allocate a new classroom $d + 1$.

Schedule lecture j in classroom $d + 1$.

$d \leftarrow d + 1$.

RETURN schedule.

Lemma.

The earliest-start-time-first algorithm can be implemented in $O(n \log n)$ time.

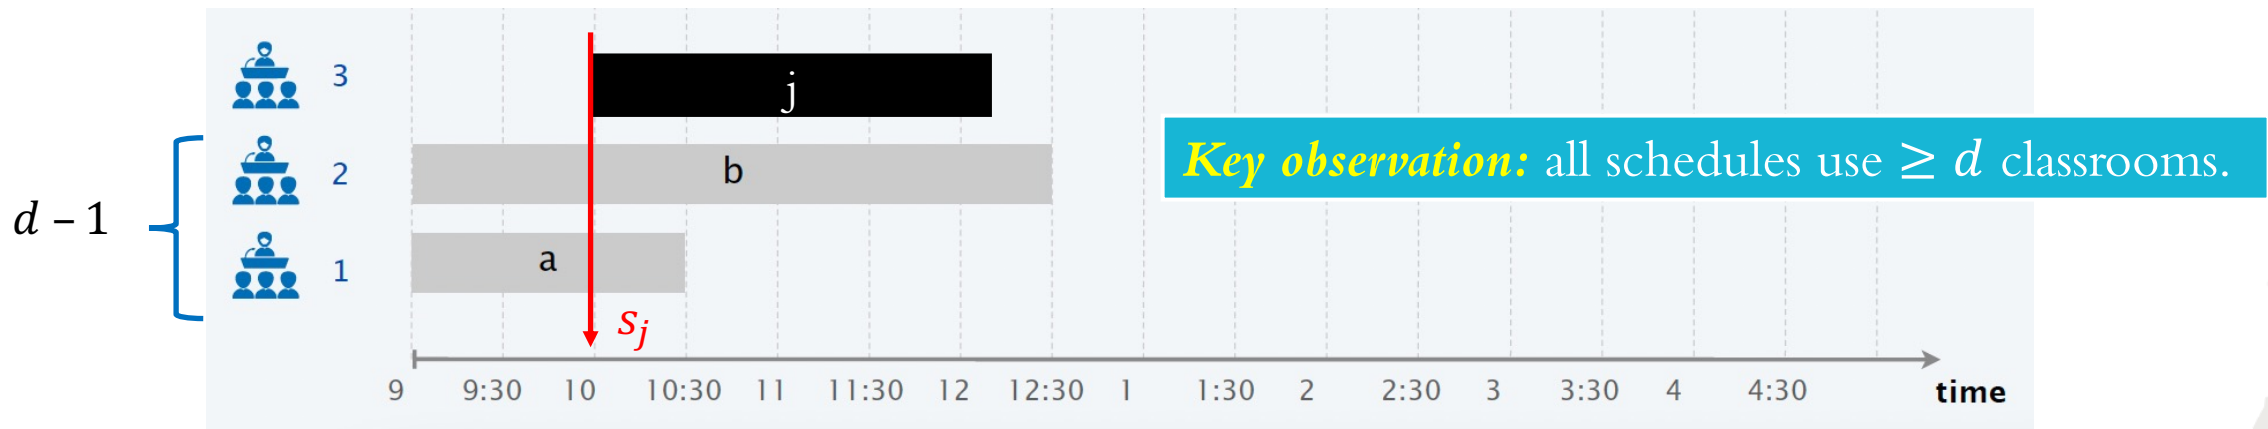
Lemma.

The earliest-start-time first algorithm never schedules two incompatible lectures in the same classroom.

INTERVAL PARTITIONING: EARLIEST-START-TIME-FIRST ALGORITHM

Theorem. Earliest-start-time-first algorithm uses #depth rooms and thus is optimal.

- Let d = number of classrooms that the algorithm allocates.
- Classroom d is opened because we needed to schedule a lecture, say j , that is incompatible with a lecture in each of $d - 1$ other classrooms.
- Thus, these d lectures each end after s_j . → The d lectures are incompatible.
- Since we sorted by start time, each of these incompatible lectures start no later than s_j . ■



SCHEDULING TO MINIMIZING LATENESS

SCHEDULING TO MINIMIZING LATENESS


Single resource processes one job at a time.

➤ Job j requires t_j units of processing time and is due at time d_j .

➤ If j starts at time s_j , it finishes at time $f_j = s_j + t_j$.

➤ **Lateness**: $l_j = \max\{0, f_j - d_j\}$.


Goal: schedule all jobs to minimize **maximum** lateness $L = \max_j l_j$.



	1	2	3	4	5	6
t_j	3	2	1	4	3	2
d_j	6	8	9	9	14	15



SCHEDULING TO MINIMIZING LATENCY



SCHEDULING TO MINIMIZING LATENESS

EARLIEST-DEADLINE-FIRST ($n, t_1, t_2, \dots, t_n, d_1, d_2, \dots, d_n$)

SORT jobs by due times and renumber so that $d_1 \leq d_2 \leq \dots \leq d_n$.

$t \leftarrow 0$.


FOR $j = 1$ **TO** n ← Process the ordered jobs one by one (immediately)

Assign job j to interval $[t, t + t_j]$.

$s_j \leftarrow t$; $f_j \leftarrow t + t_j$.

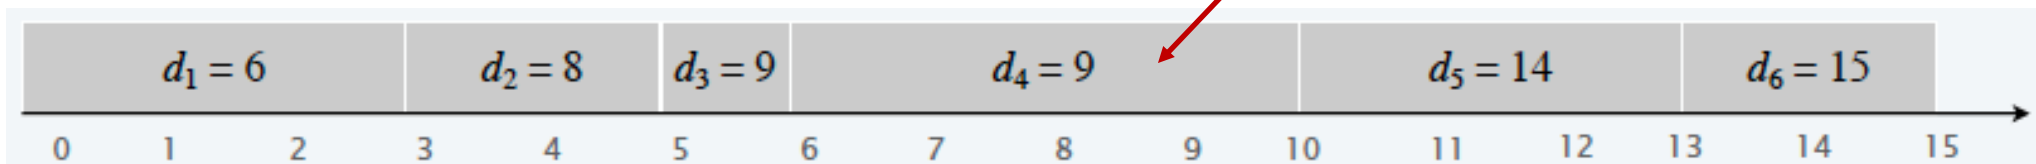
$t \leftarrow t + t_j$.

RETURN intervals $[s_1, f_1], [s_2, f_2], \dots, [s_n, f_n]$.



	1	2	3	4	5	6
t_j	3	2	1	4	3	2
d_j	6	8	9	9	14	15

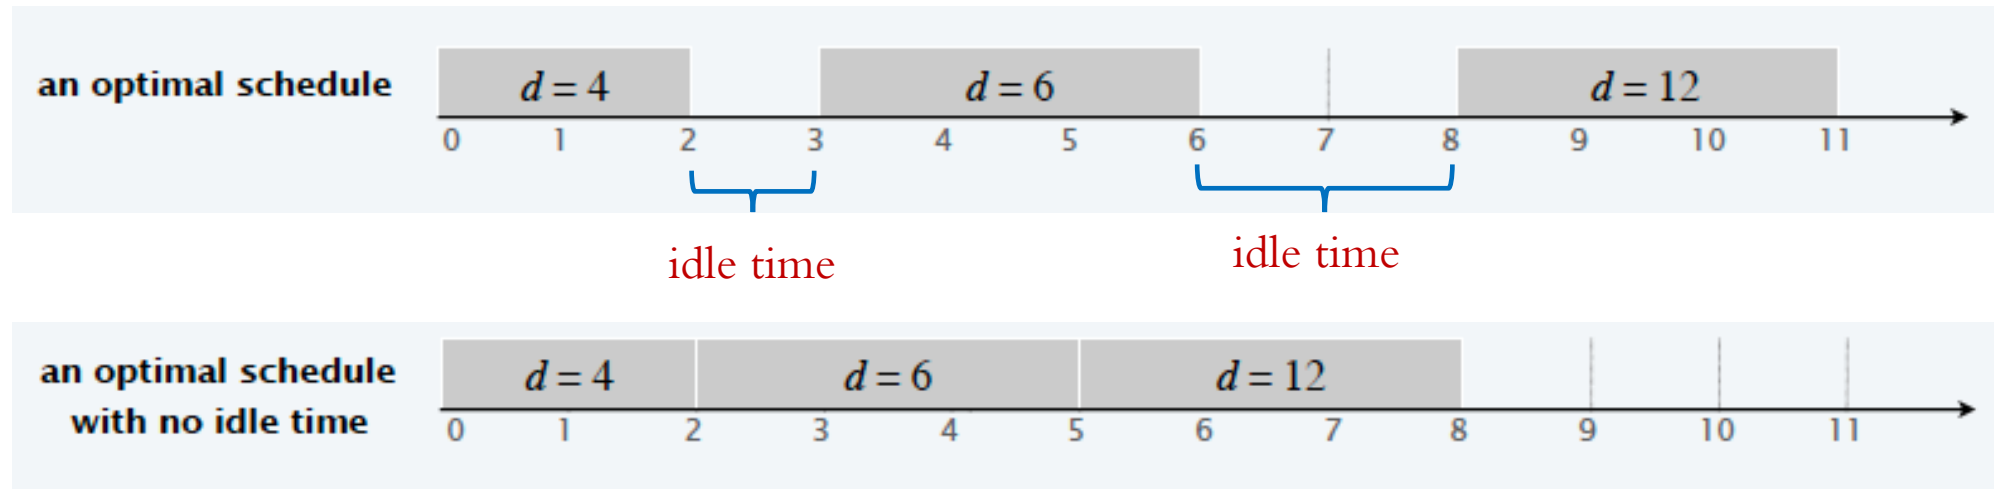
$l_4 = 1$



SCHEDULING TO MINIMIZING LATENESS

Properties for optimal schedules.

Observation 1. There exists an optimal schedule with no idle time.

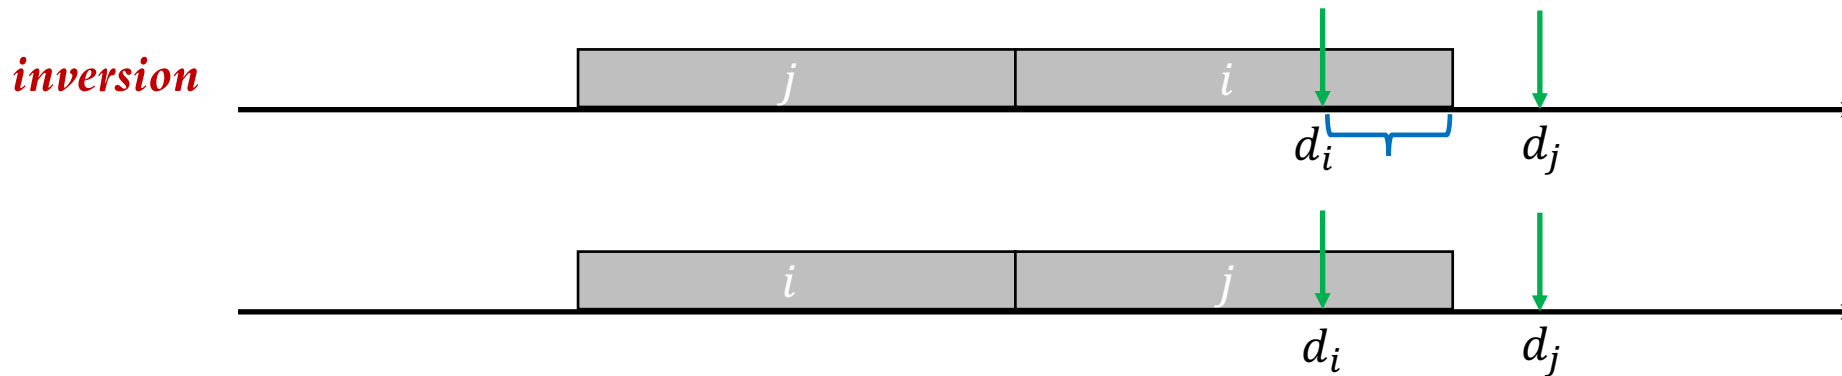


Observation 2. The earliest-deadline-first schedule has no idle time.

SCHEDULING TO MINIMIZING LATENESS

Definition. Given a schedule S , an **inversion** is a pair of jobs i and j such that: $d_i < d_j$ but j is scheduled before i .

or $i < j$ for ordered jobs



swap makes the schedule better!

Observation 3. The earliest-deadline-first schedule is the **unique** idle-free schedule with no inversions.

SCHEDULING TO MINIMIZING LATENESS

Observation 4. If an idle-free schedule has an inversion, then it has an **adjacent inversion**.

two inverted jobs scheduled consecutively

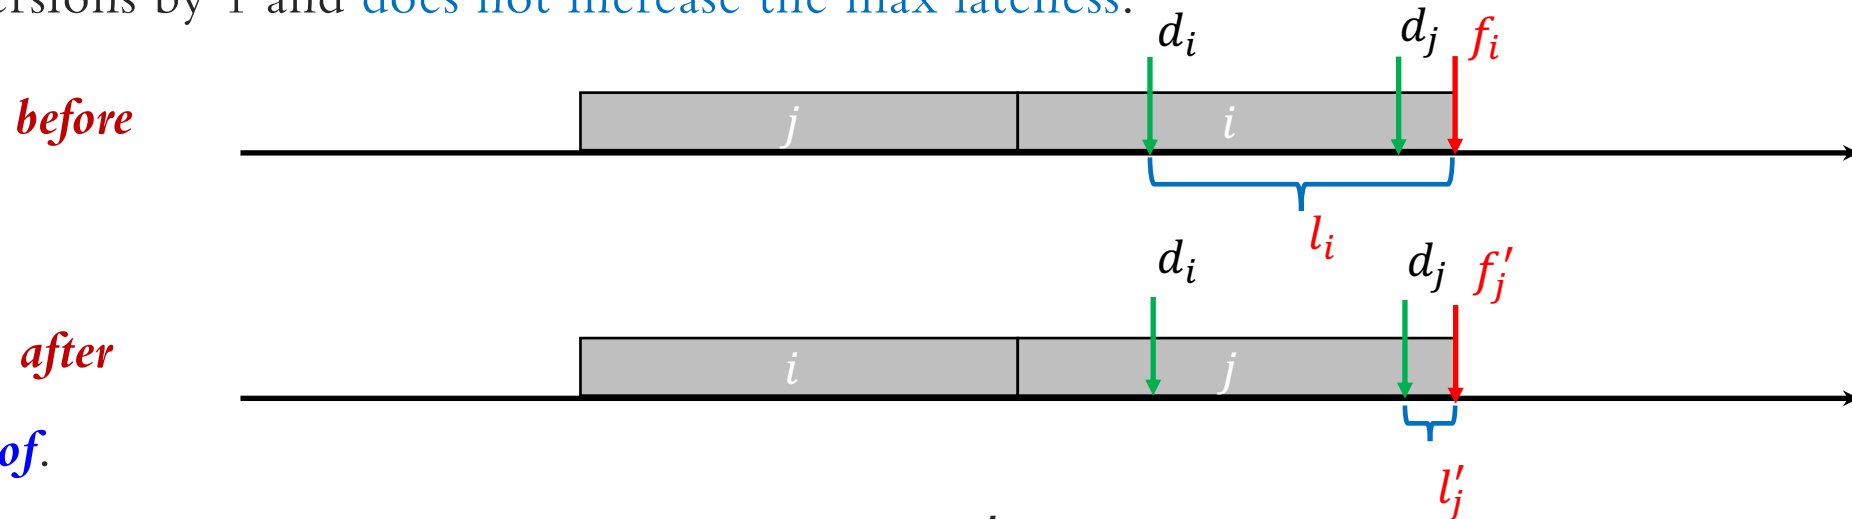
Proof.

- Let $i-j$ be a **closest** inversion. $d_j > d_i$
- Let k be element immediately to the right of j .
 - **Case 1:** $d_j > d_k$. Then $j-k$ is an adjacent inversion.
 - **Case 2.** $d_j < d_k$. Then $i-k$ is a closer inversion. ■



SCHEDULING TO MINIMIZING LATENESS

Key Claim. Exchanging two **adjacent**, **inverted** jobs i and j reduces the number of inversions by 1 and **does not increase the max lateness**.



Proof.

- Let l be the lateness before the swap, and let l' be it afterwards.
- $l'_k = l_k$ for all $k \neq i, j$.
- $l'_i \leq l_i$
- If job j is late, $l'_j = f'_j - d_j = f_i - d_j \leq f_i - d_i \leq l_i$. ■

SCHEDULING TO MINIMIZING LATENESS

Theorem. The earliest-deadline-first schedule S is optimal.

Proof. [by contradiction]

- Define S^* to be an optimal schedule with the **fewest inversions**.
- Can assume S^* has no idle time. \longrightarrow **Observation 1**
- **Case 1:** S^* has no inversions. Then $S = S^*$. \longrightarrow **Observation 3**
- **Case 2:** S^* has an inversion.
 - Let $i - j$ be an **adjacent** inversion \longrightarrow **Observation 4**
 - Exchanging jobs i and j decreases the number of inversions by 1 without increasing the max lateness \longrightarrow **Key Claim**
 - Contradicts “**fewest inversions**” part of the definition of S^* . ■

GREEDY ANALYSIS STRATEGIES

Greedy algorithm stays ahead.

- Show that after each step of the greedy algorithm, its solution is **at least as good as any other algorithm's**.
- [Interval scheduling]

Structural.

- Discover a simple “structural” bound asserting that **every possible solution must have a certain value**. Then show that your algorithm always achieves this bound.
- [Interval partitioning]

Exchange argument.

- Gradually **transform any solution to the one found by the greedy algorithm** without hurting its quality.
- [Minimizing lateness, Interval scheduling]

Thank You!